



Hands-On Material:

github.com/db-tu-dresden/SIGMOD25-OptimizerTutorial

Reproducible Prototyping of Query Optimizer Components

Rico Bergmann and Dirk Habich

Technische Universität Dresden, Germany



SIGMOD'25 Tutorial, June 27, 2025, Berlin, Germany

Presenters

Rico Bergmann



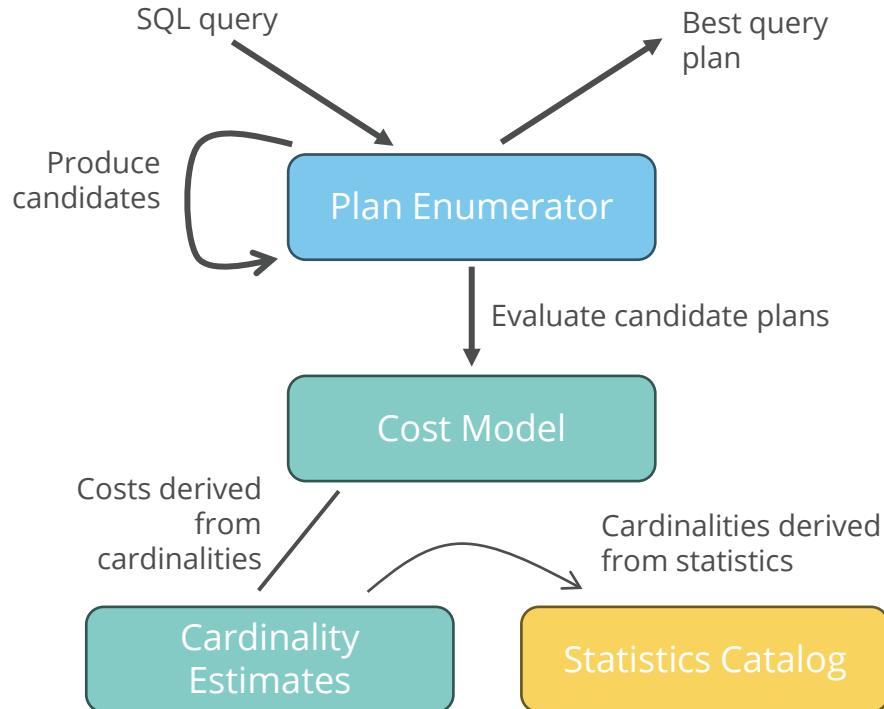
- PhD student (2nd year)
- Technische Universität Dresden, Germany
- Research focus
 - Combining “traditional” and learned query optimization
 - Research tools for query optimization

Dirk Habich



- Associate Professor
- Technische Universität Dresden, Germany
- Research focus
 - Database systems on modern hardware
 - Query optimization
 - Habilitation in 2018 on „In-Memory Database Query Processing on Large-Scale Multiprocessor Systems“

Textbook Query Optimizer



Current Research – Cost Model

Zero-Shot Cost Models for Out-of-the-box Learned Cost Prediction

Johannes Härtig*

Tobias Hirschmann
Technische Universität Darmstadt

Sebastian Schuster
Microsoft Research Asia

ABSTRACT

In this paper, we propose a novel cost model, which handles cost estimation that generalizes to unseen instances. It leverages a learned cost estimator that is trained on a large dataset to predict a range of training times for every new state transition. This allows us to estimate the cost of a sequence of transitions without having to train a separate model for each transition. To enable such zero-shot predictions, we propose a novel learning paradigm that uses a learned cost estimator as a teacher to guide the student to learn the cost of a specific transition. We evaluate our approach on a learned cost model for the Microsoft Dynamics AX system. Our learned cost model is able to predict the cost of a sequence of transitions with a relative error of less than 10% for 90% of the sequences. Furthermore, we show that our learned cost model is able to predict the cost of a sequence of transitions with a relative error of less than 10% for 90% of the sequences. Furthermore, we show that our learned cost model is able to predict the cost of a sequence of transitions with a relative error of less than 10% for 90% of the sequences.

* Denotes equal contribution. Correspondence to: Tobias Hirschmann (tobias.hirschmann@tu-darmstadt.de).

Copyright © 2019, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

AAAI-19-0773.pdf

<div data-bbox="1

Current Research



Join Ordering

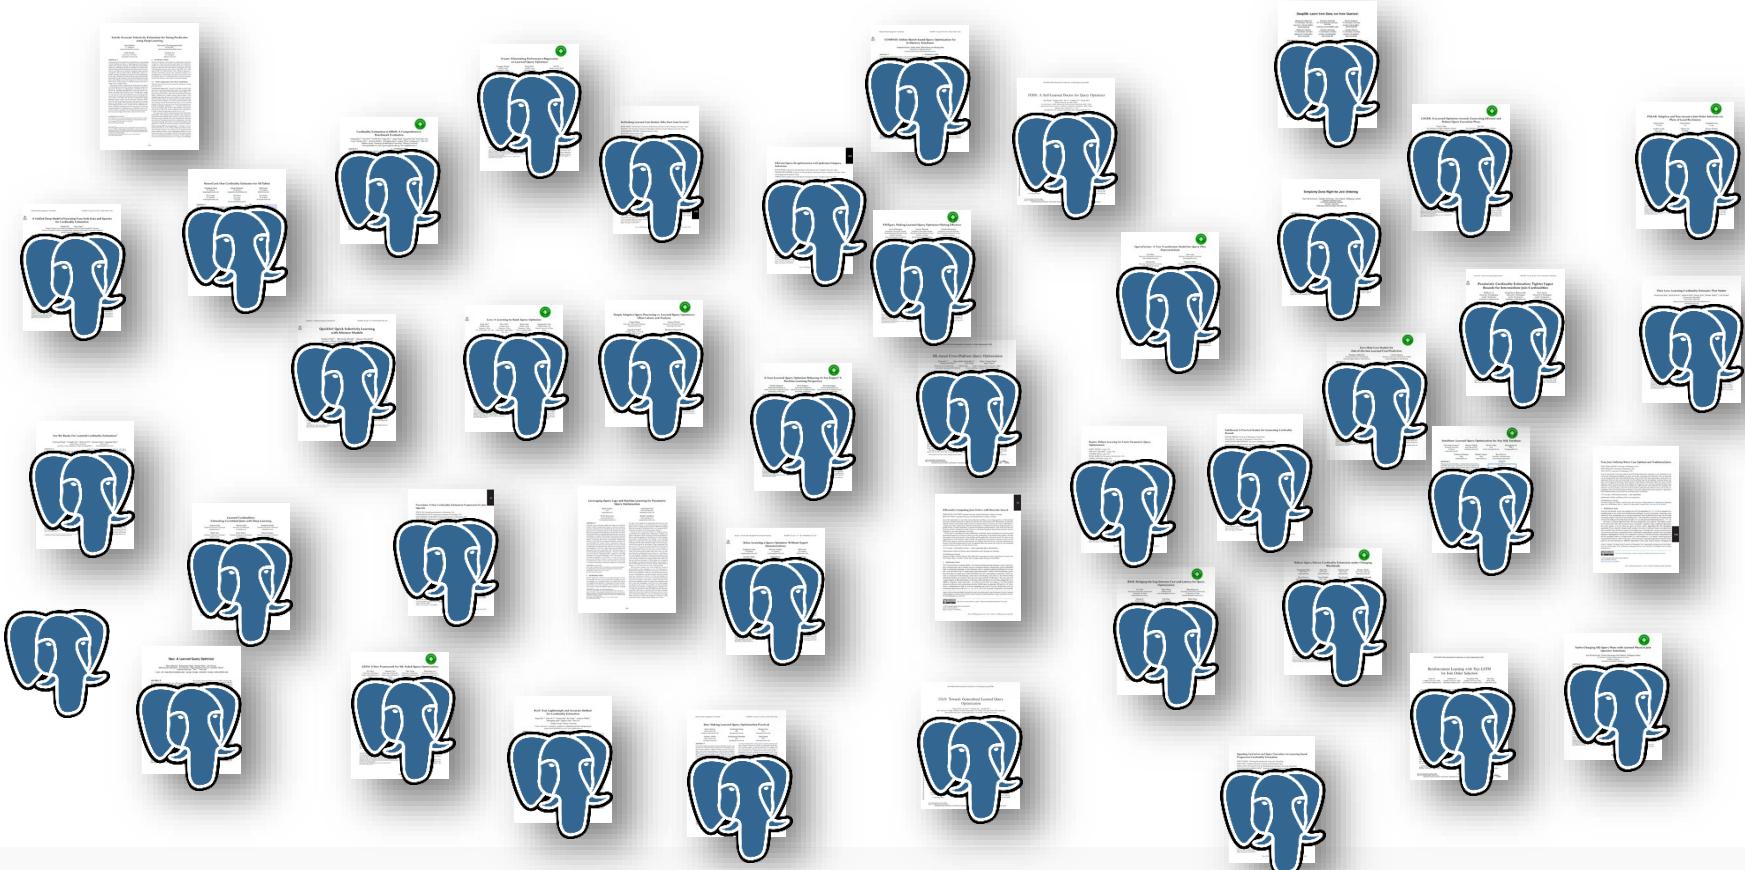
Operator Selection

Pessimistic Optimization

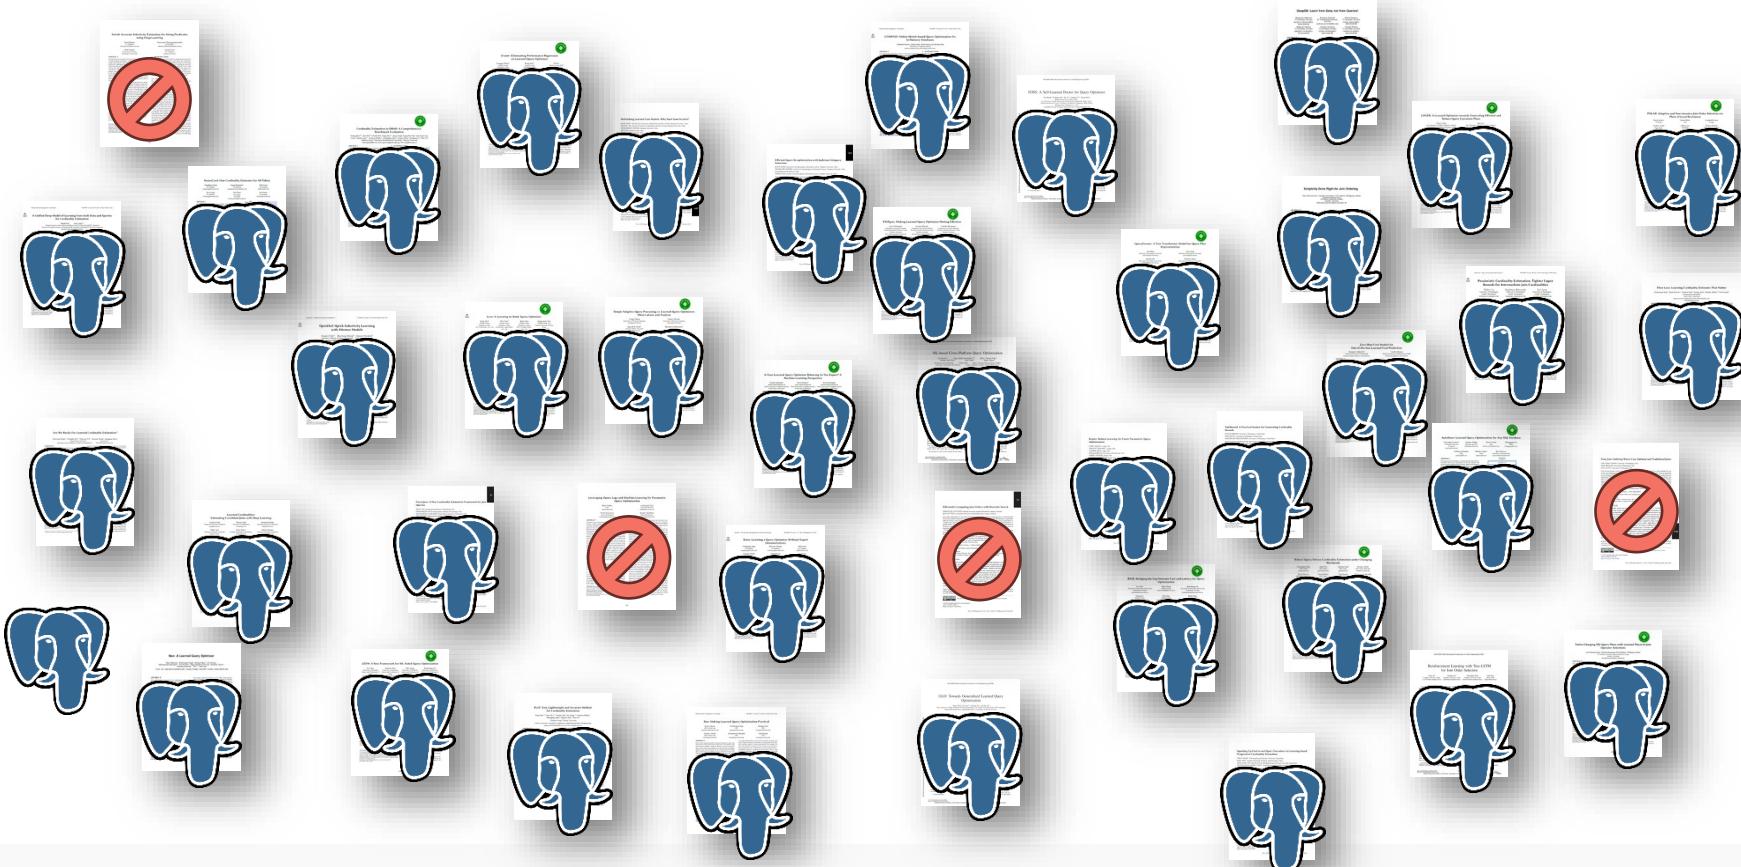
Current Research – Common Ground



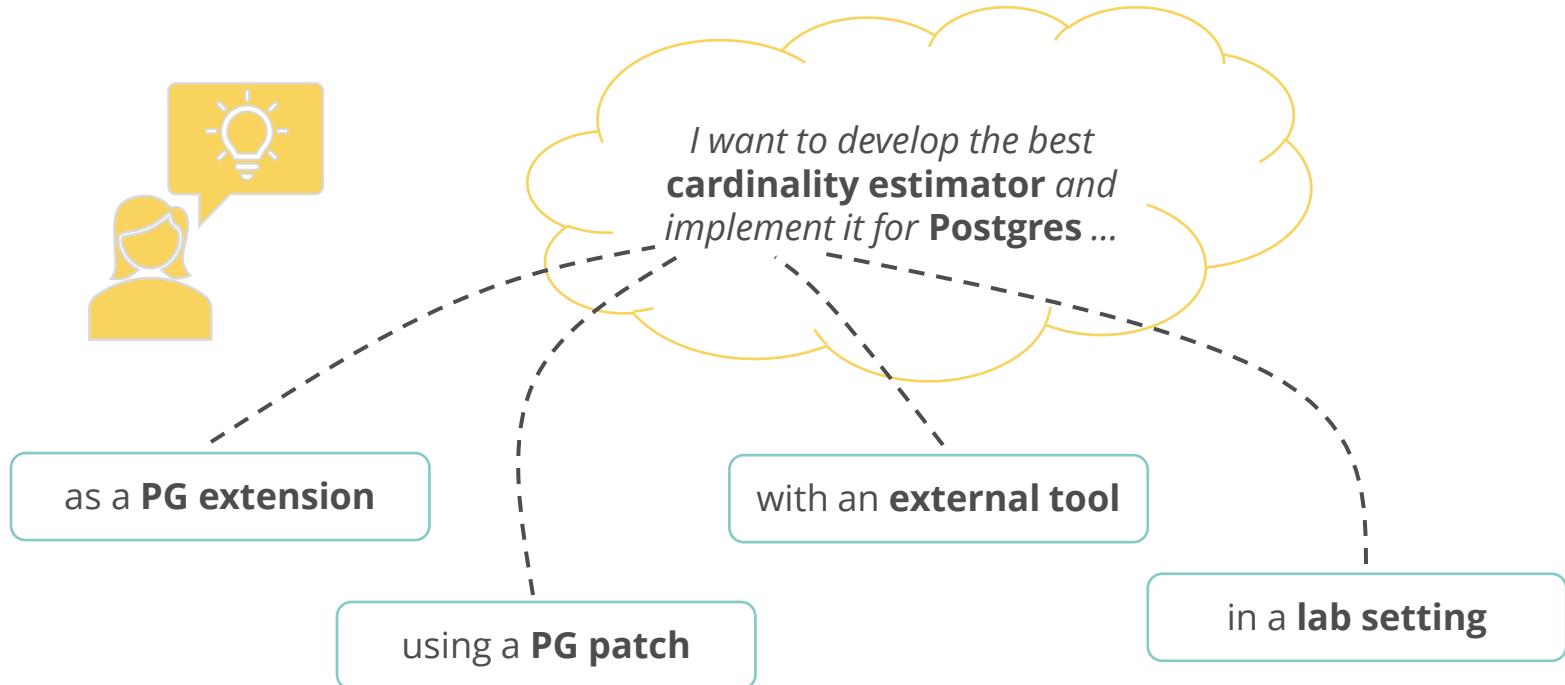
Current Research – Common Ground



Current Research – Common Ground



From Idea To Implementation



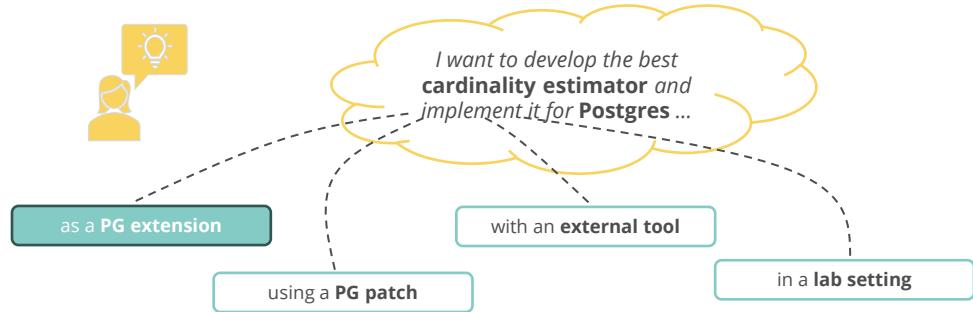
Implementation

Postgres Extension

- Postgres includes powerful extension mechanism
- Dynamically added to the server process
- Provides abstract functions to plug in custom functionality

Existing Hooks

- Complete optimization
 - planner_hook
- Join ordering + operator selection
 - join_search_hook
- Operator selection
 - set_rel_pathlist_hook
 - set_join_pathlist_hook
 - create_upper_paths_hook



Downsides

- Relies on available hooks
 - E.g., currently no hook for cardinality estimator or cost model
- Requires deep understanding of PG internals

```
double max_column_frequency(PlannerInfo *root, RelOptInfo *rel, Var *column) {
    AttStatsSlot sslot;
    VariableStatData vardata;
    double max_freq;
    examine_variable(root, (Node *) column, 0, &vardata);
    get_attstatsslot(&sslot, vardata.statsTuple, STATISTIC_KIND_MCV, InvalidOid,
                      ATTSTATSSLOT_VALUES | ATTSTATSSLOT_NUMBERS);

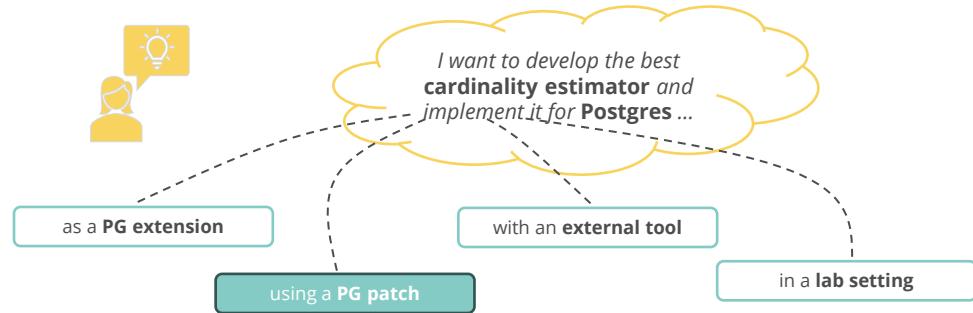
    if (sslot.nvalues > 0)
        /* We found an MCV list, use it */
        max_freq = sslot.numbers[0] * rel->tuples; /* index 0 is the highest frequency */
    else {
        /* No MCV found, assume equal distribution */
        bool default_est; /* ignored, keep compiler quiet */
        double ndistinct = get_variable_numdistinct(&vardata, &default_est);
        if (ndistinct < 0) /* negative means ND is a fraction of the table size */
            ndistinct *= -1.0 * rel->tuples;
        max_freq = rel->tuples / ndistinct;
    }

    ReleaseVariableStats(vardata);
    return max_freq;
}
```

Implementation

Postgres Patch

- Direct modifications of the PG source code
 - Enables changes to query optimizer and execution engine
- Adaptive query optimization becomes possible



```
8347  /*  
8348   * clamp_row_est  
8349 @@ -4733,6 +4758,10 @@ set_joinrel_size_estimates(PlannerInfo *root, RelOptInfo *rel,  
8350  
8351  
8352  
8353     if (enable_lero) {  
8354       lero_pgsql_set_joinrel_size_estimates(root, rel, outer_rel,  
8355                                         inner_rel, sjinfo, restrictlist);  
8356     }  
8357   }  
8358  
8359  /*
```

Downsides

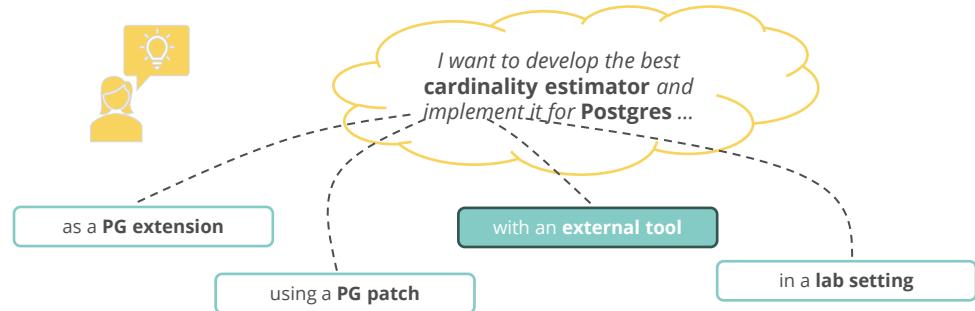
- Requires deep understanding of PG internals
- Changes are hard to track
- Specific to (minor) releases of PG

Implementation

pg_hint_plan

- Extension to embed optimizer decisions in comments
- Hints can be generated through arbitrary software
- Allows good tooling integration (e.g., ML frameworks)

```
/*+ NestLoop(t mi) IndexScan(mi) */
SELECT count(*)
FROM title t, movie_info mi
WHERE t.id = mi.movie_id
    AND t.production_year > 2010
```



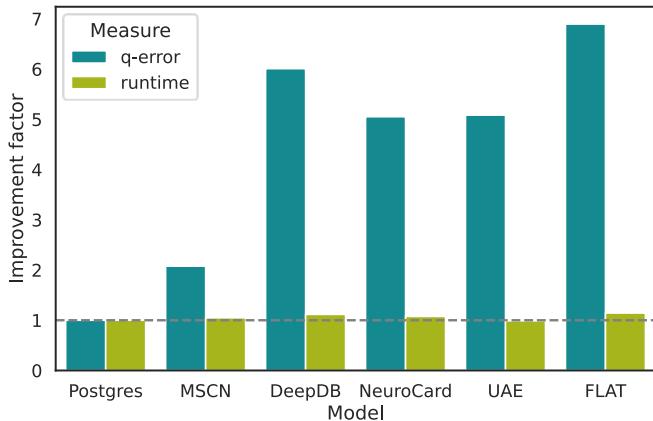
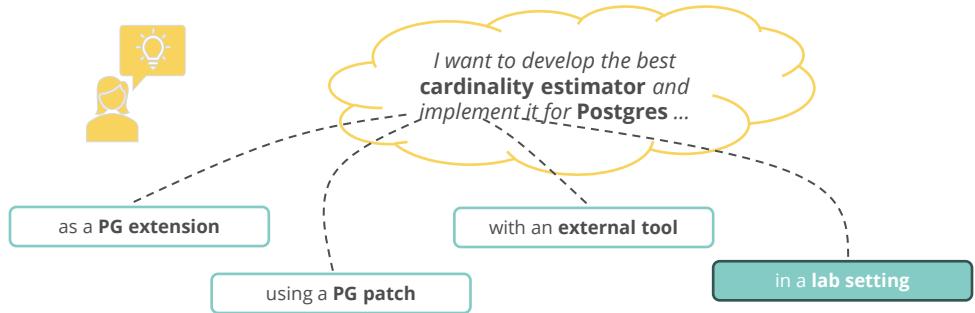
Downsides

- Opaque implementation – no standardized generation of hints
- Limited to the available hints
 - No cardinalities for base tables
 - No parallel workers

Implementation

Artificial Setting

- Don't implement for a real database
- Instead, use a "laboratory setting"
- E.g., measure q-error instead of runtime

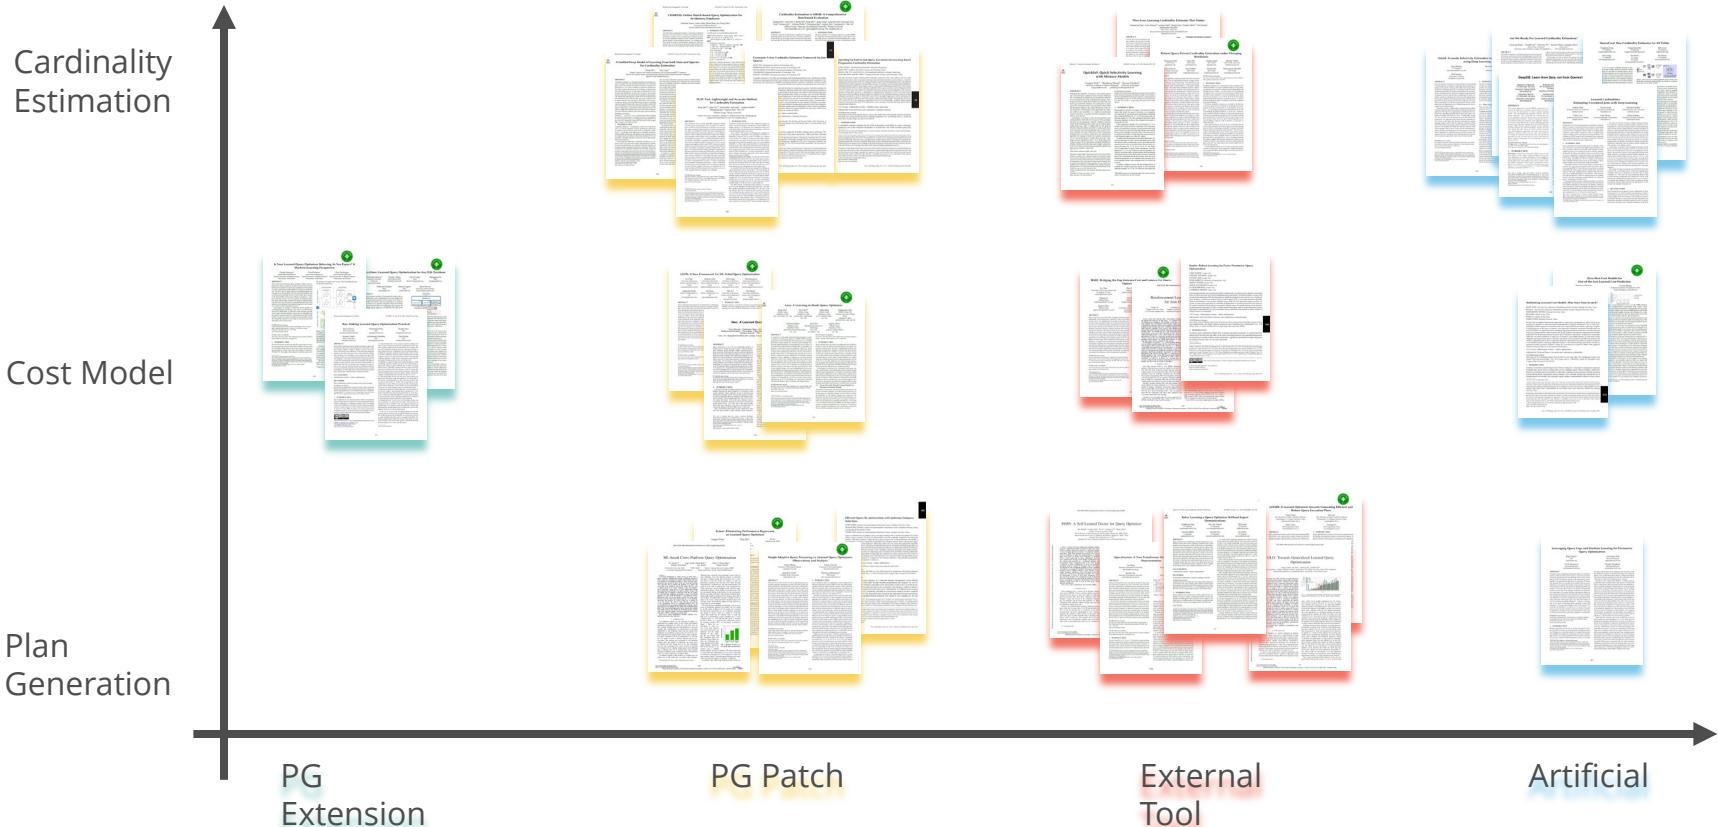


Downsides

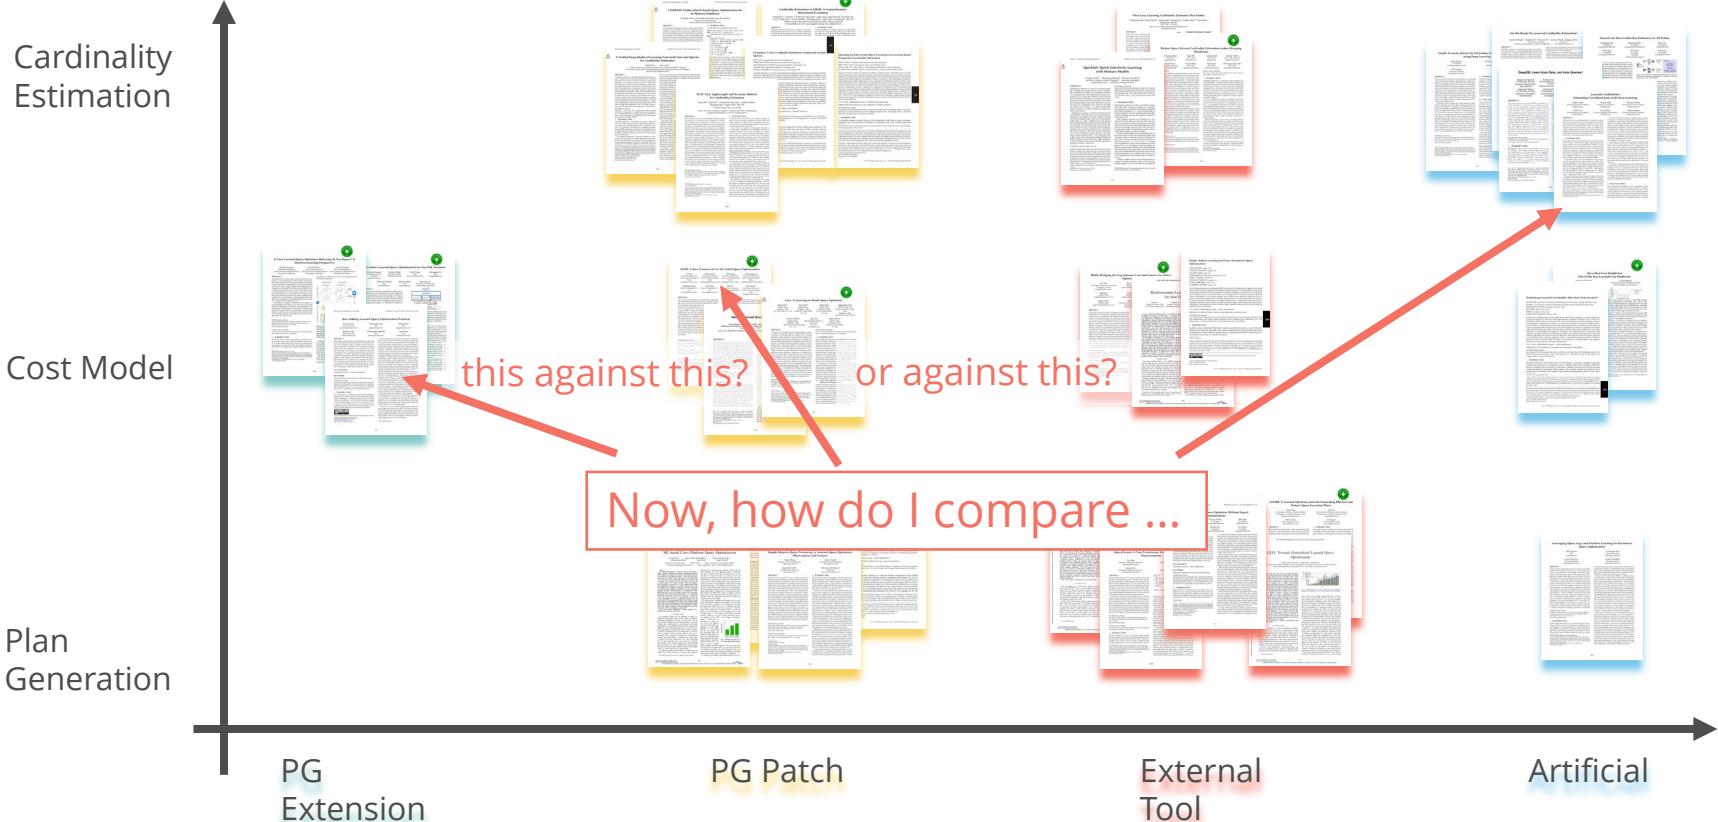
- No real-world execution – how to ensure correlation?

Numbers based on Han et al.:
Cardinality Estimation in DBMS: A Comprehensive Benchmark Evaluation (VLDB'21)

Current Research Revisited



Current Research Revisited



Let's compare with BAO (Marcus et al.)

BAO in a nutshell

- Uses search space restrictions to obtain different query plans
- Estimates the plan quality via a learned model
- Selects the (estimated) best plan for execution
- SIGMOD best paper award 2021

Implementation details

- Implemented as a PG extension
- Communicates with a Python server
- Based on PG v12.4

Marcus et al.: "BAO: Making Learned Query Optimization Practical" (SIGMOD'2021)

SIGMOD '21, June 20–25, 2021, Virtual Event, China

Research Data Management Track Paper

Bao: Making Learned Query Optimization Practical

Ryan Marcus
MIT & Intel Labs
rymanmarcus@mit.edu
Parimkar Negi
MIT
pnegi@mit.edu
Hongzhi Mao
MIT
hongzhi@mit.edu

Nesime Tatbul
MIT & Intel Labs
tatbul@csail.mit.edu
Mohammad Alizadeh
MIT
alizadeh@csail.mit.edu
Tim Kraska
MIT
kraska@csail.mit.edu

ABSTRACT
Recent efforts applying machine learning techniques to query optimization have shown promising results. However, these approaches often suffer from overhead, inability to adapt to changes, and poor tail performance. In this paper, we propose Bao, a learned query optimizer that can make probabilistically expensive operation in practice (thus why we wish to call it "learned") practical. Bao takes advantage of the window built-in existing query optimizers by providing per-query optimization hints. Bao uses a novel sampling, a well-studied reinforcement learning algorithm. As opposed to previous work, Bao does not require any prior knowledge about the system, sampling, or workload. Bao is able to handle schema changes in query workloads, data, and schemas. Experimentally, we show that Bao outperforms state-of-the-art query optimizers in terms of end-to-end query execution performance, including tail latency for several workloads containing long-running queries. In short, Bao is able to provide better performance compared with a commercial system.

CCS CONCEPTS
Information systems → Query optimization.

KEYWORDS
Database management, machine learning, reinforcement learning

ACM Reference Format:
Ryan Marcus, Parimkar Negi, Hongzhi Mao, Nesime Tatbul, Mohammad Alizadeh, and Tim Kraska. 2021. Bao: Making Learned Query Optimization Practical. In Proceedings of the 2021 International Conference on Management of Data (SIGMOD '21), June 20–25, 2021, Virtual Event, China, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3453009.3453038>

1 INTRODUCTION
Query optimization is an important task for database management systems (DBMS). It is a complex problem that requires elements of query optimization – cardinality estimation and cost modeling – to work correctly [38]. Several works have applied machine learning (ML) to query optimization [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54]. While all of these new solutions demonstrate some promise, they also raise the question that many of the techniques are yet practical, as they suffer from several fundamental problems:

- (1) Long training time. Most proposed machine learning techniques require a large amount of training data to learn, which may have a positive impact on query performance. For example, ML-powered cardinality estimators based on supervised learning require millions of training examples to learn [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54].
- (2) Inability to adjust to data and workload changes. While proposed reinforcement learning approaches are able to handle impractical changes in query workload, data, or schema, they can make them very slow. This is because the learned models need to be retrained when data changes, or risk becoming stale.
- (3) Inability to provide a way for database administrators to influence the learned models.
- (4) Integrations costs. To the best of our knowledge, all previous learned optimizers are still research prototypes, offering little to no integration with real-world DBMSs. They are usually limited to a subset of standard SQL, not to mention vendor-specific features. Hence, fully integrating learned optimizers into a production DBMS is a non-trivial task. Moreover, the fact that learned optimizers are not yet fully integrated into a DBMS makes it difficult to implement them in a production environment.

To the best of our knowledge, Bao (Bao) is the first learned query optimizer that addresses all of the above-mentioned problems. Bao is fully integrated into PostgreSQL, as an extension, and can be used by any PostgreSQL user. The database administrator (DBA) just needs to download our open-source module¹, and even has the option to selectively turn the learned optimizer on or off for specific queries.

1. https://bao.csail.mit.edu

This work is licensed under a Creative Commons Attribution International 4.0 License
© 2021. Copyright held by the owner/author(s).
<https://doi.org/10.1145/3453009.3453038>

1279

Let's compare with BAO (Marcus et al.)

```
~/pg-17/contrib/bao/pg_extension [master] $ make && make install
gcc -Wall -Wmissing-prototypes -Wpointer-arith -Wdeclaration-after-statement -Werror=vla
-Wendif-labels -Wmissing-format-attribute -Wimplicit-fallthrough=3 -Wcast-function-type
-Wshadow=compatible-local -Wformat-security -fno-strict-aliasing -fwrapv -fexcess-
precision=standard -Wno-format-truncation -Wno-stringop-truncation -O2 -fPIC -
fvisibility=hidden -I. -I./ -I/home/rico/pg-17/build/include/postgresql/server -
I/home/rico/pg-17/build/include/postgresql/internal      -D_GNU_SOURCE    -c -o main.o
main.c
In file included from main.c:9:
bao_bufferstate.h:13:10: fatal error: utils/reffilenodemap.h: No such file or directory
 13 | #include "utils/reffilenodemap.h"
   | ^~~~~~
compilation terminated.
make: *** [<>builtin>: main.o] Error 1
```

We can't reproduce on PG 17!

OK, what about Lero? (Zhu et al.)

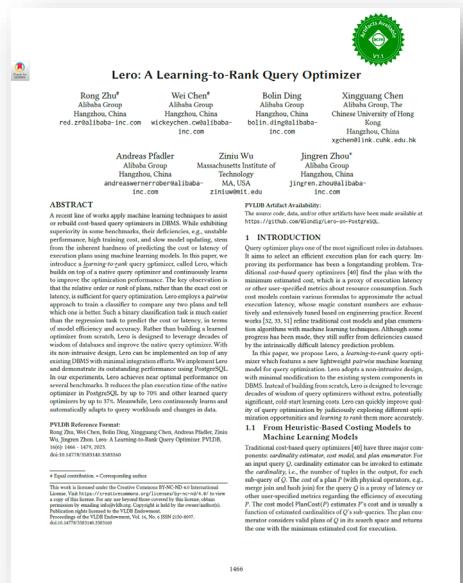


Lero in a nutshell

- Uses modified cardinality estimates to obtain different query plans
 - Ranks the plans using a BAO-style model
 - Selects the (ranked) best plan for execution

Implementation details

- Implemented as a PG patch
 - Communicates with a Python server
 - Based on PG v13.1



Zhu et al.: "I ero: A Learning-to-Rank Query Optimizer" (VLDB'2023)

OK, what about Lero? (Zhu et al.)

```
~/pg-17 [master] $ git apply contrib/lero/0001-init-lero.patch
error: patch failed: src/backend/Makefile:19
error: src/backend/Makefile: patch does not apply
error: patch failed: src/backend/optimizer/path/costsize.c:107
error: src/backend/optimizer/path/costsize.c: patch does not apply
error: patch failed: src/backend/optimizer/plan/planner.c:64
error: src/backend/optimizer/plan/planner.c: patch does not apply
error: patch failed: src/backend/utils/misc/guc.c:98
error: src/backend/utils/misc/guc.c: patch does not apply
error: patch failed: src/include/optimizer/cost.h:67
error: src/include/optimizer/cost.h: patch does not apply
```

We can't reproduce on PG 17!

Can we at least use MSCN? (Kipf et al.)



Artificial evaluation

=
No implementation on database system

=

We can't reproduce on PG 17!

Kipf et al.: "Learned Cardinalities: Estimating Correlated Joins with Deep Learning" (CIDR'2019)

Learned Cardinalities: Estimating Correlated Joins with Deep Learning

Andreas Kipf
Technical University of Munich
kipf@in.tum.de

Viktor Lek
Technische Universität München
leojin@in.tum.de

Peter Boncz
Centrum Wiskunde & Informatica
boncze@wmi.nl

Thomas Kipf
University of Amsterdam
takipf@science.ru.nl

Bernhard Radke
Technical University of Munich
radke@in.tum.de

Alfred Kemper
Technische Universität München
kemper@in.tum.de

We argue that machine learning is a highly promising technique for solving the cardinality estimation problem. Learning can be used to estimate the cardinality of correlated joins by learning query features and the output being the estimated cardinality. In contrast to sampling-based cardinality estimation techniques, which are often limited to simple join types like index structures [15] and join ordering [23], the current techniques based on basic per-table statistics are not very good for large joins. However, we show that our model does not have to be perfect, it just needs to be better than the current, inaccurate sampling-based approaches. We show that our machine learning model can directly be leveraged by existing sophisticated query optimizers without requiring any other changes to the database system.

INTRODUCTION
Query optimization is fundamentally based on cardinality estimation. To be able to choose between different plan alternatives, the optimizer needs to estimate the cost of each plan, which includes the estimated result sizes. It is well known, however, that the estimates produced by all widely-used database systems are routinely wrong [1, 2, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23]. The biggest challenge in cardinality estimation are correlated joins. Correlated joins are common in movie recommendation systems such as Netflix [24] and in movie databases such as Movie Database (IMDb). French actors are more likely to participate in romantic movies than actors of other genres. This is an open area of research. One state-of-the-art proposal in this area is Index-based cardinality estimation [10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23].

There are no qualifying samples to start with (i.e., under selective queries there are no samples available) and therefore the optimizer has to fall back to sampling-based cardinality estimation. Sampling-based cardinality estimation is a well-known technique for estimating the size of sets [25, 26, 27]. It involves drawing random samples from the set and then counting the number of samples that belong to the set. This is achieved by using a binomial distribution to estimate the probability of a sample belonging to the set [25].

Deep learning has been applied to query optimization by three recent papers [13, 20, 27] that formulate join ordering as a reinforcement learning problem and use ML to find query plans. This work, in contrast, focuses on the cardinality estimation problem in isolation. This focus is motivated by the fact that modern join estimation algorithms can find the optimal join order for queries with dozens of relations [28]. Cardinality estimation, on the other

Could frameworks help us?

Ideally, we would like:



Broad support for different optimization strategies



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

Existing Frameworks

Research Frameworks

Apache Calcite

- Java-based framework to implement entire optimizer
- Community-driven development
- Deep integration in Java/Apache ecosystem
- Allows customization of a traditional optimizer model

<https://calcite.apache.org/>



	Broad support for different optimization strategies
	Transparent implementation of prototypes
	Execution on Postgres
	Easy comparison of different prototypes

Research Frameworks

LingoDB

<https://www.lingo-db.com/>

- Research system focused on MLIR-based query processing
- Developed at TUM
- Optimizer implemented as LLVM compiler passes



	Broad support for different optimization strategies
	Transparent implementation of prototypes
	Execution on Postgres
	Easy comparison of different prototypes

Research Frameworks

PilotScope

- Research framework focused on ML4DB
- Not limited to query optimization
 - E.g., also support for index recommendation
- Developed at Alibaba
- Uses Push/Pull connectors on an actual DBMS

<https://github.com/alibaba/pilotscope>



	Broad support for different optimization strategies
	Transparent implementation of prototypes
	Execution on Postgres
	Easy comparison of different prototypes

PostBOUND

Overview

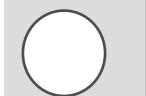
- Python framework for research in query optimization
- Two primary use cases:
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



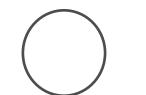
Broad support for different optimization strategies



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes





Overview

- Python framework for research in query optimization
- Two primary use cases
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimization strategies



Broad support for different optimization strategies



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

- Provides different optimization pipelines
- Pipelines correspond to different optimizer architectures
- Researchers select suitable pipeline
- Researchers implement interfaces of the pipeline



Overview

- Python framework for research in query optimization
- Two primary use cases
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



Broad support for different optimization approaches



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

- Use hinting to enforce PostBOUND's optimization decisions
- Hints overwrite optimizer of the target database



Overview

- Python framework for research in query optimization
- Two primary use cases
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



Broad support for different optimization strategies



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

- Benchmark “support suite”, e.g.,
- Query plan analysis
- Automatic cache prewarming
- Pre-defined workloads + data sets (IMDB, Stats, Stack, SSB)

Non-functional goals

- Easy to use
- Nice programming experience
 - Powerful abstractions
 - Low boilerplate
 - High customizability
- Fast onboarding
 - Provide many optimization algorithms out of the box
- Low configuration overhead
 - Use sensible defaults whenever possible
- Good documentation





PostBOUND

Non-functional requirements

- Easy to use
- Nice API
- Performance
- Local storage
- Hybrid storage
- Fast
- Parallel
- Optimization
- Low overhead
- Usage
- Good documentation



PostBOUND 0.16.0

Search Ctrl + K

10 Minutes to PostBOUND

Setup

Core concepts

Advanced usage

Cookbook

API

<https://postbound.readthedocs.io/en/latest/>



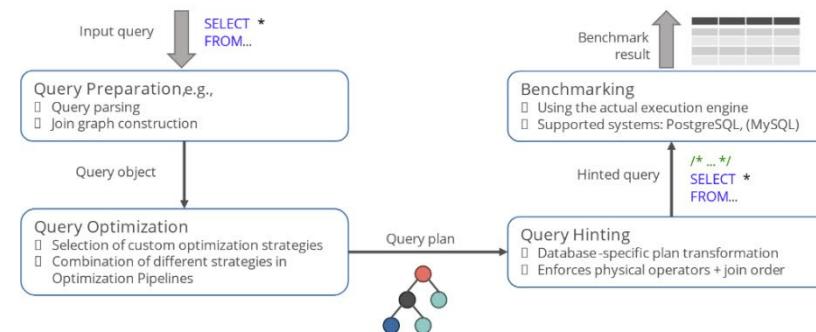
PostBOUND documentation

PostBOUND is a Python framework for research in query optimization. It provides a high-level interface to implement novel optimization algorithms and to analyze their performance.

Tip

New users should start by reading the [10 Minutes to PostBOUND](#) tutorial and the [Setup](#) guide. The remainder of the documentation describes the different parts of the framework in more detail.

In general, using the PostBOUND framework for optimizer research uses a workflow similar to the following:



Typical workflow for using PostBOUND for optimizer research.

Contents

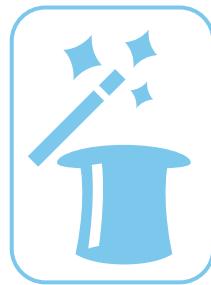
- PostBOUND documentation
- Contents
- Example
- History
- Citation
- Indices and tables
- References



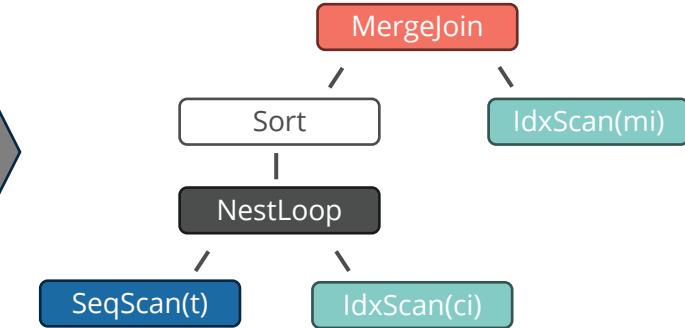
Plan Enforcement via Query Hinting



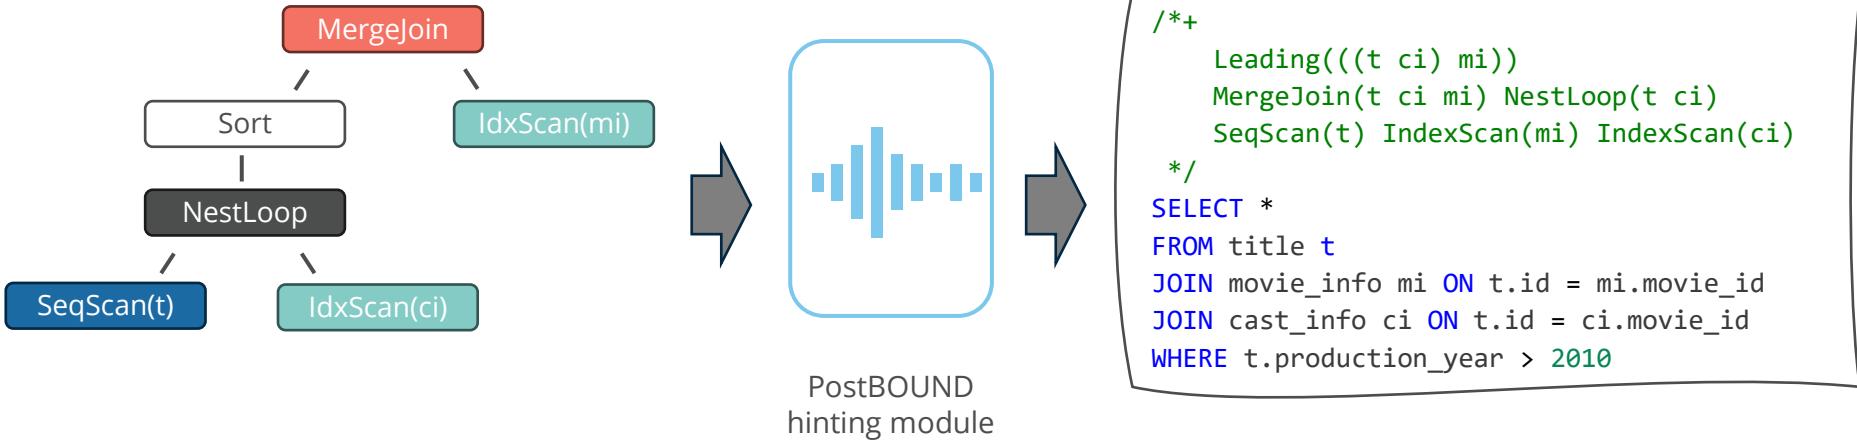
```
SELECT *  
FROM title t  
JOIN movie_info mi ON t.id = mi.movie_id  
JOIN cast_info ci ON t.id = ci.movie_id  
WHERE t.production_year > 2010
```



PostBOUND “magic”
(for now)



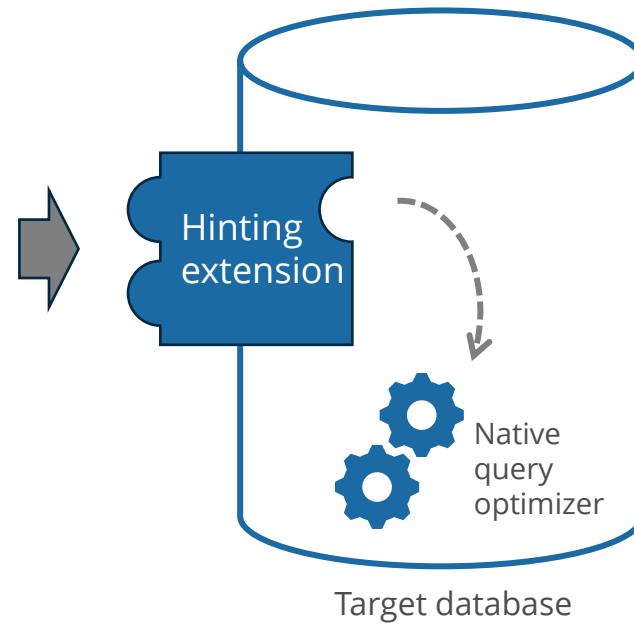
Plan Enforcement via Query Hinting



Plan Enforcement via Query Hinting



```
/*+
  Leading(((t ci) mi))
  MergeJoin(t ci mi) NestLoop(t ci)
  SeqScan(t) IndexScan(mi) IndexScan(ci)
*/
SELECT *
FROM title t
JOIN movie_info mi ON t.id = mi.movie_id
JOIN cast_info ci ON t.id = ci.movie_id
WHERE t.production_year > 2010
```

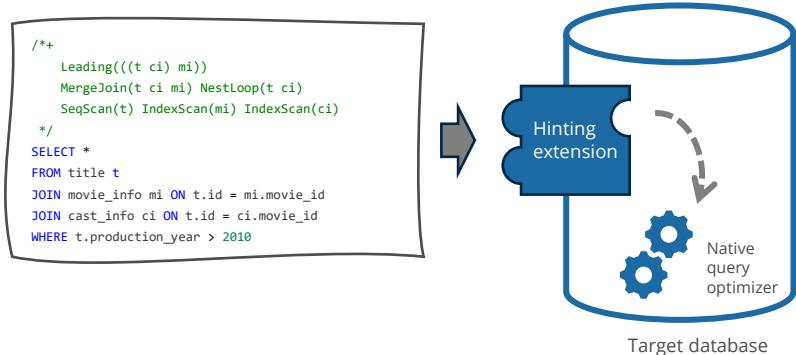


Plan Enforcement via Query Hinting



Supported Hinting Backends

- pg_hint_plan
 - Hints for join order
 - Hints for most physical operators
 - Hints for join cardinalities
- pg_lab
 - Same as pg_hint_plan, but additionally:
 - Hints for parallel plans
 - Hints for base table cardinalities



Consequences

- Can only enforce plan features that are supported by the hinting backend
- Currently no support for re-optimization approaches

"Postgres for Optimizer Research"

- Lightweight fork of PostgreSQL
- Adds additional extension points to the optimization workflow
- Provides fine-grained hinting capabilities

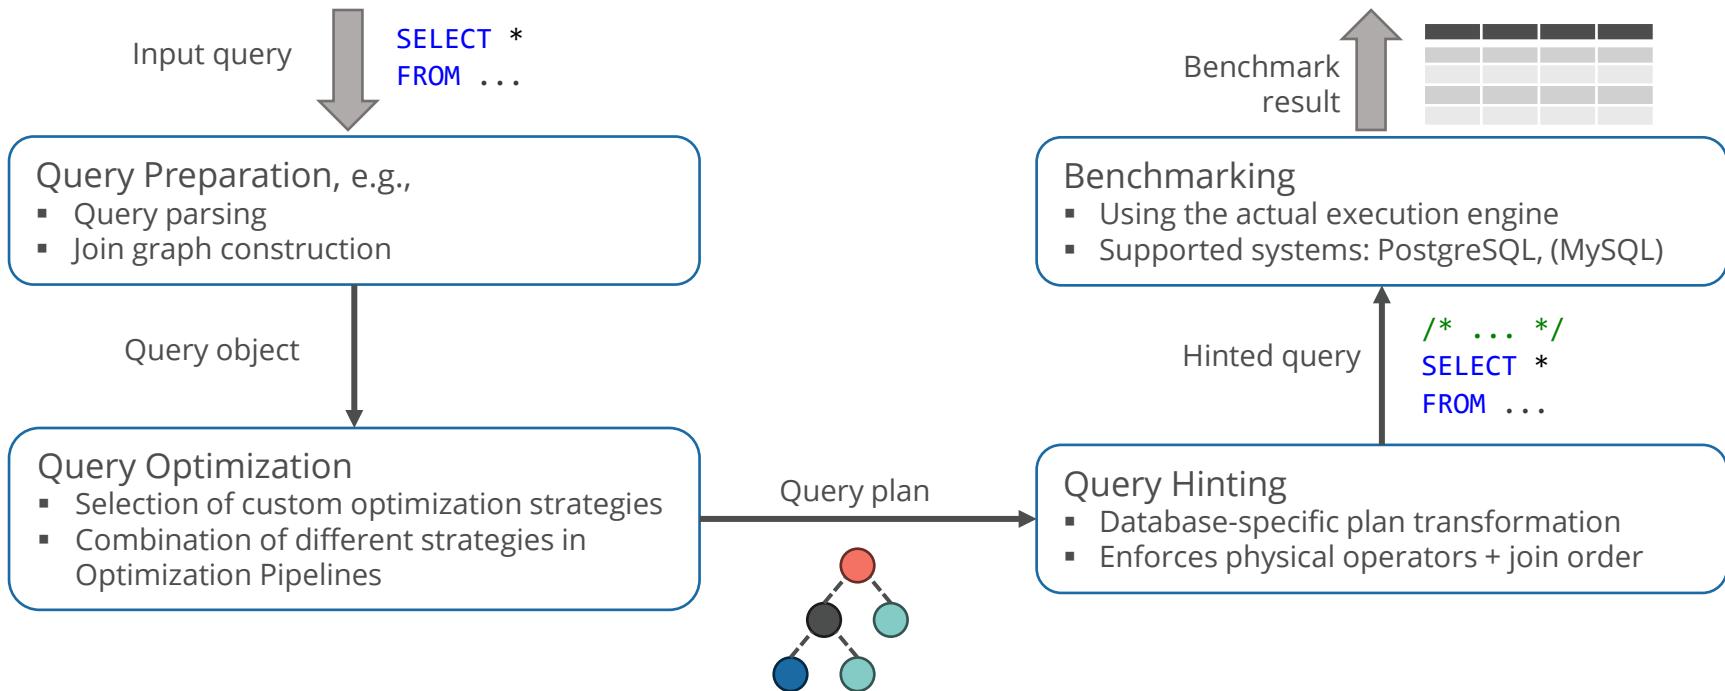


Extension points

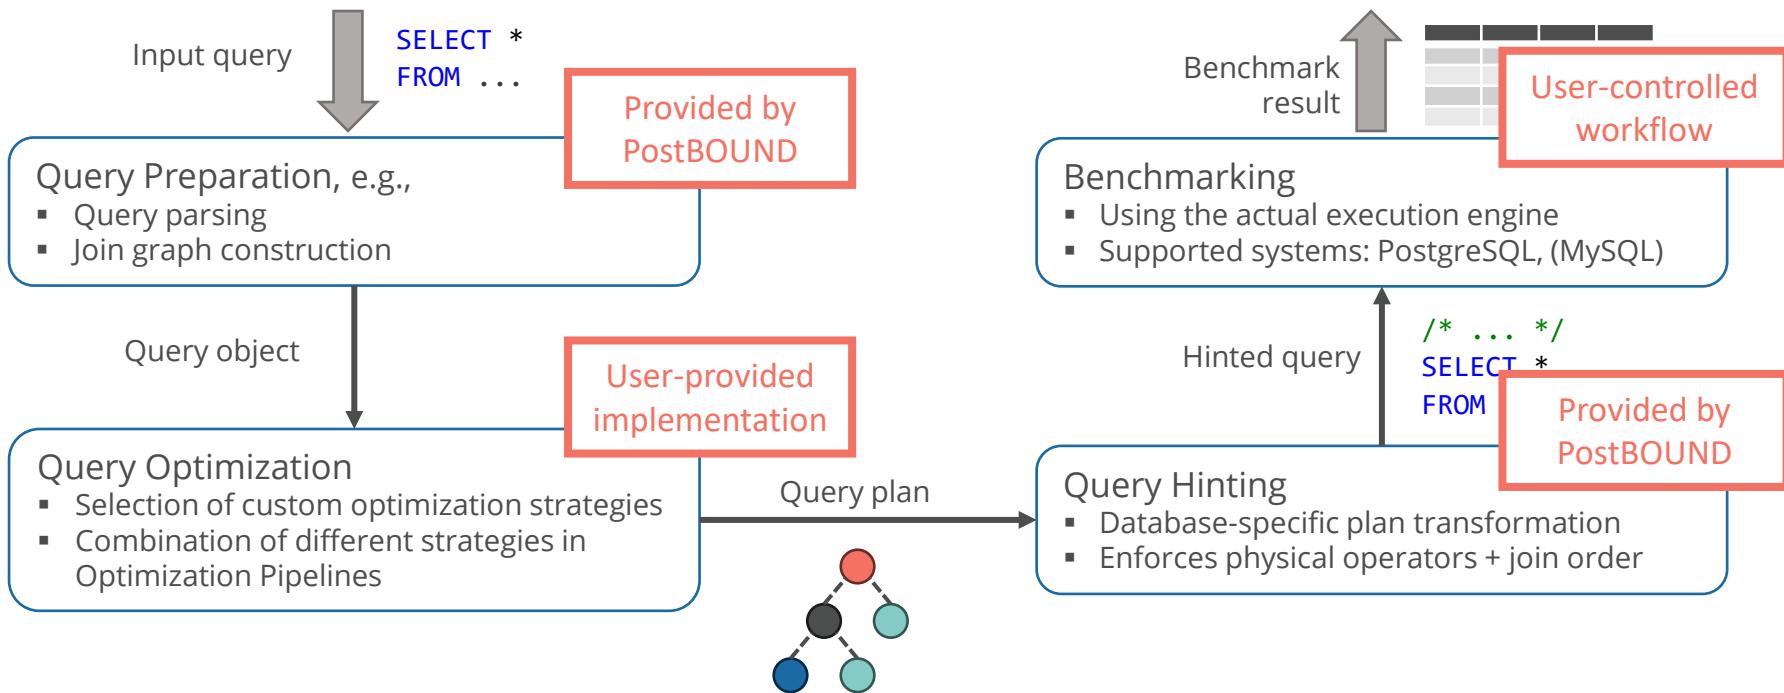
- Cardinality estimation
- Cost estimation
- Path pruning
- Parallel worker estimation



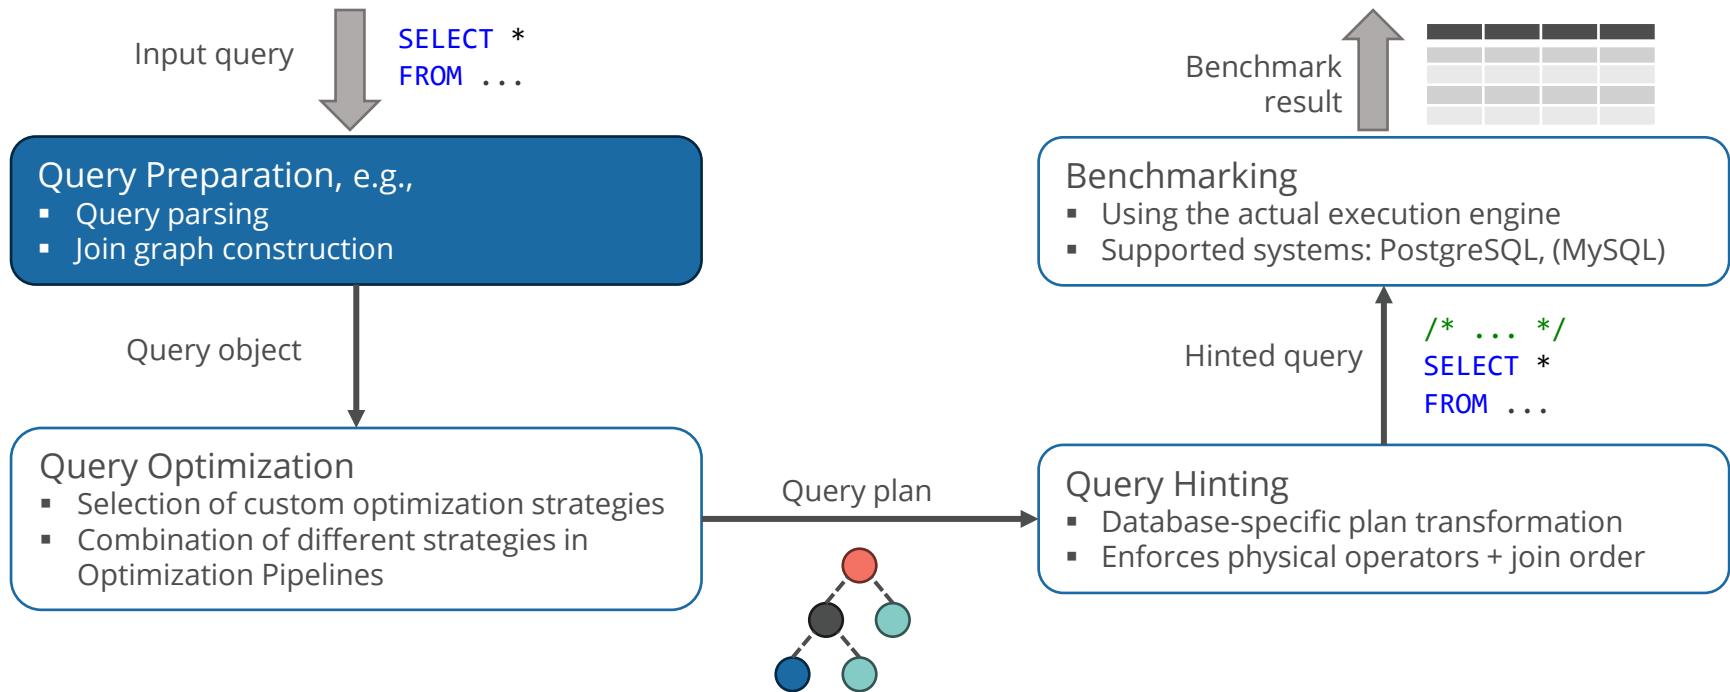
PostBOUND Workflow



PostBOUND Workflow



PostBOUND Workflow



Stats Benchmark

Dataset

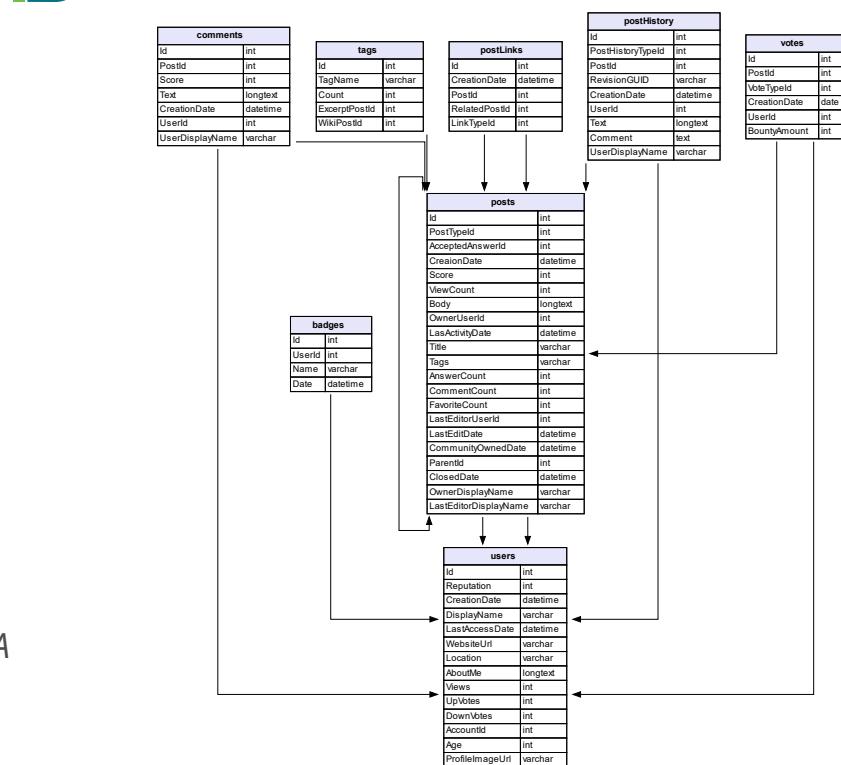
- Based on the Statistics Stackoverflow ([CrossValidated](#))
- Models relation between users, their questions, and answers
- 8 tables with complex (foreign-key) correlations

Workload

- 146 queries in total, between 1 and 6 joins
- Auto-generated queries based on complex templates

Purpose

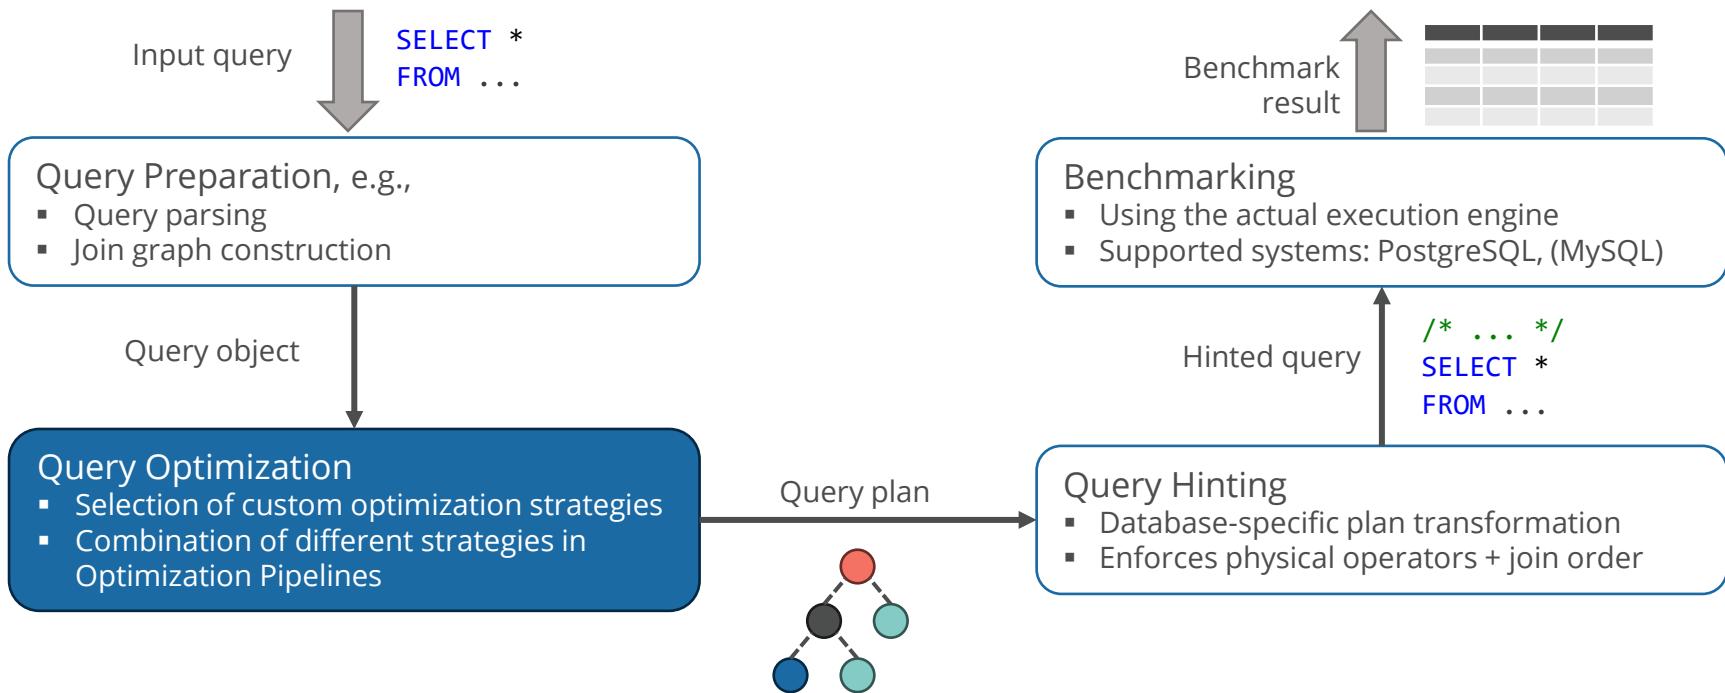
- Lightweight benchmark (<125MB compressed) for cardinality estimation
- Challenging correlations, skewed data distributions
- Presented by Han et al.: “*Cardinality Estimation in DBMS: A Comprehensive Benchmark Evaluation*” (VLDB 2022, [Link](#))





Demo

PostBOUND Workflow



Optimizer Implementation

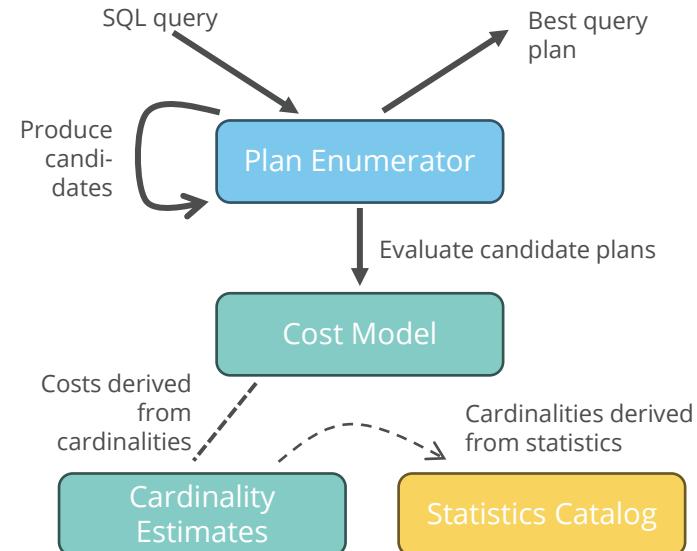


Optimization pipelines

- Mental models for different optimizer architectures
- Pipelines define specific optimization steps (e.g., cardinality estimation)
- Researchers need to implement (subsets of) the stages
- “Fill the gaps” principle
 - Researchers only need to worry about their specific stages
 - PostBOUND uses reasonable defaults for the rest

Textbook pipeline

- Available stages:
 - Plan enumerator
 - Cost model
 - Cardinality estimator
- Defaults: native estimators and simulated enumerator



Optimizer Implementation

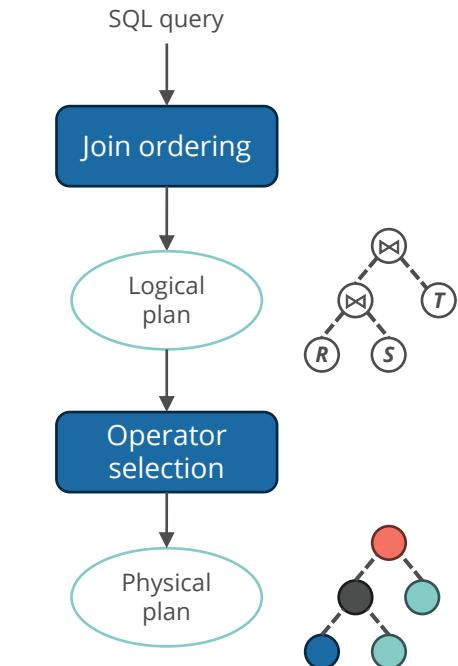


Optimization pipelines

- Mental models for different optimizer architectures
- Pipelines define specific optimization steps (e.g., cardinality estimation)
- Researchers need to implement (subsets of) the stages
- “Fill the gaps” principle
 - Researchers only need to worry about their specific stages
 - PostBOUND uses reasonable defaults for the rest

Multi-stage pipeline

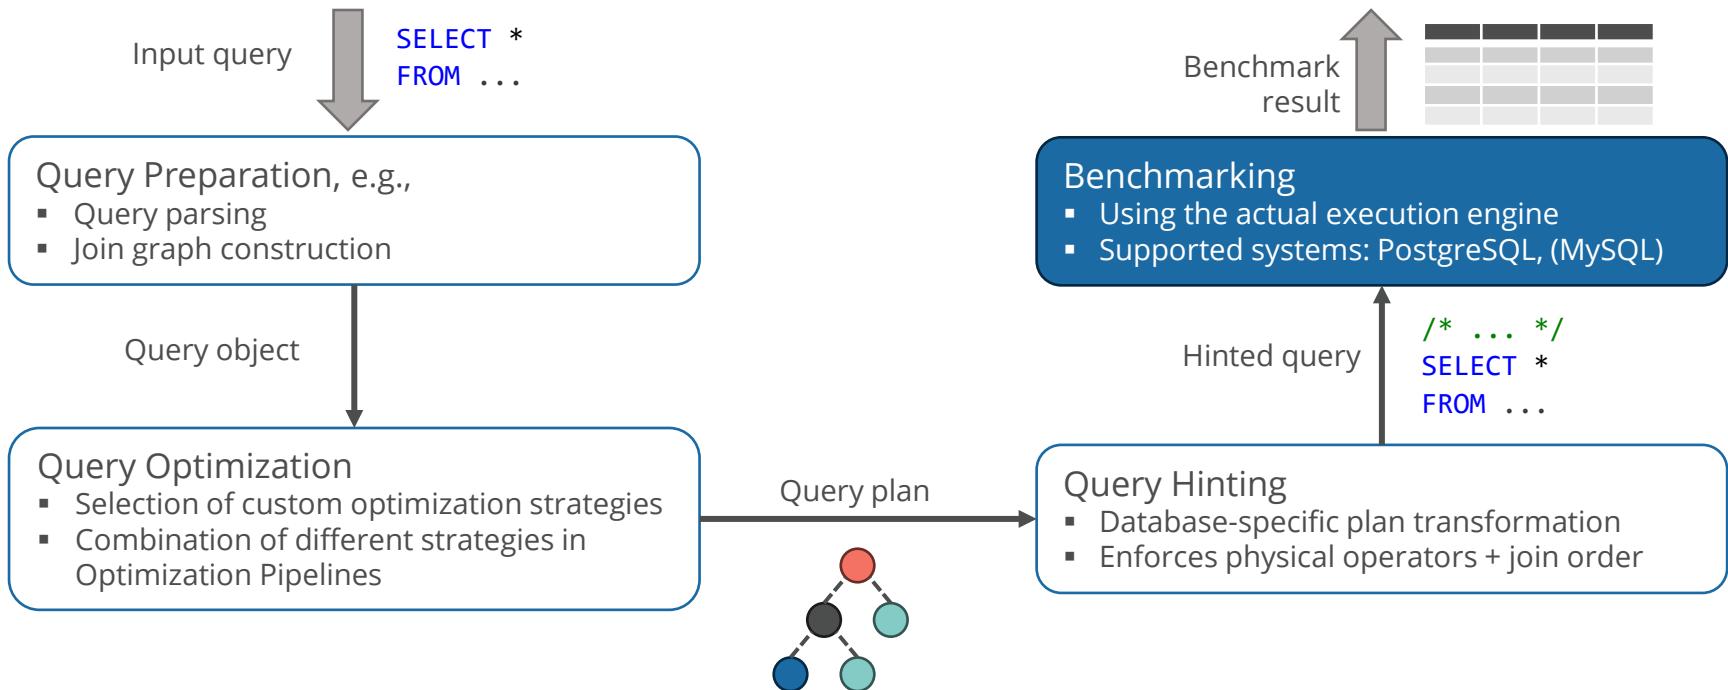
- Available stages:
 - Join ordering
 - Operator selection
 - (Plan parameterization)
- Defaults: native strategies





Demo

PostBOUND Workflow



Benchmarking



Benchmark execution

- Goal: reproducible + transparent evaluation
- Query preparation for common tasks (e.g., warm start)
- Tools generate many additional data points
 - Pipeline configuration
 - Database schema
 - ...
- Integrated into Python data analysis ecosystem (e.g., pandas)



Support Tooling

- Setup scripts for common databases + workloads (IMDb, SSB, Stats, StackOverflow)
- Setup scripts for different Postgres versions





Dresden
Database
Research Group

Demo



Implementing Your Own Optimizer

Hands-On Material:
github.com/db-tu-dresden/SIGMOD25-OptimizerTutorial



More Examples

SIGMOD Research Paper

- „An Elephant Under The Microscope“
- Broad analysis of the interaction of optimizer components
- Custom optimization “strategies” for
 - Join ordering
 - Cardinality estimation
 - Pessimistic query optimization



An Elephant Under The Microscope: Analyzing The Interaction Of Optimizer Components In PostgreSQL

RICO BERGMANN, Technische Universität Dresden, Germany

CLAUDIO HARTMANN, Technische Universität Dresden, Germany

DIRK HABICH, Technische Universität Dresden, Germany

WOLFGANG LEHNER, Technische Universität Dresden, Germany

Despite an ever-growing corpus of novel query optimization strategies, the interaction of the core components of query optimizers is still not well understood. This situation can be problematic for two main reasons: On the one hand, this may cause surprising results when two components influence each other in an unexpected way. On the other hand, this can lead to wasted effort in regard to both engineering and research, e.g., when two implemented optimizer components are dwarfed by the performance of one by far more efficient component. Therefore, we argue that analyzing the interactions between components in a single optimizer component requires a thorough understanding of how these changes might affect the other components. To achieve this understanding, we present results of a systematic experimental analysis of the interplay in the traditional optimizer component of the widely-used PostgreSQL system as well as our own implementation. Our results can help to revisit the core building blocks of such an optimizer, i.e. per-column statistics, cardinality estimation, cost model, and plan generation. In particular, we analyze how these building blocks influence each other and how they react when faced with faulty input, such as imprecise cardinality estimates. Based on our results, we then draw conclusions and make recommendations on how these should be taken into account.

CCS Concepts: -Information systems → -Query optimization.

Additional Key Words and Phrases: Query Optimization; Data Analysis; Cardinality Estimation; Join Enumeration; Cost Model

ACM Reference Format:

Rico Bergmann, Claudio Hartmann, Dirk Habich, and Wolfgang Lehner. 2025. An Elephant Under The Microscope: Analyzing The Interaction Of Optimizer Components In PostgreSQL. *Proc. ACM Manag. Data* 3, 1 (SIGMOD), Article 9 (February 2025), 28 pages. <https://doi.org/10.1145/3799659>

1 Introduction

Although query optimization in relational database systems has been studied in decades of research [1–9, 20–23], there is still no consensus on how to do query optimization the “right” way and novel approaches are published each year. Those contributions usually focus on different aspects of the optimization process, for example by introducing new strategies for cardinality estimation [27, 63], by presenting improved cost models [23, 36], or by introducing refined strategies for plan enumeration [21, 34]. However, these improvements are typically made in isolation, without taking the interaction of the various query optimizer components into account and without comprehensively analyzing the effect of the other components. Since a query optimizer is a complex system, it is often the case that changes in one component have a significant impact on other components. For example, changes in the cost model can affect the cardinality estimation, which in turn can affect the plan enumeration. This lack of understanding can lead to suboptimal query plans and inefficient execution.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee or condition, provided that: (1) copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other use, contact the owner/authors.

© 2025 Copyright held by the owner/authors.

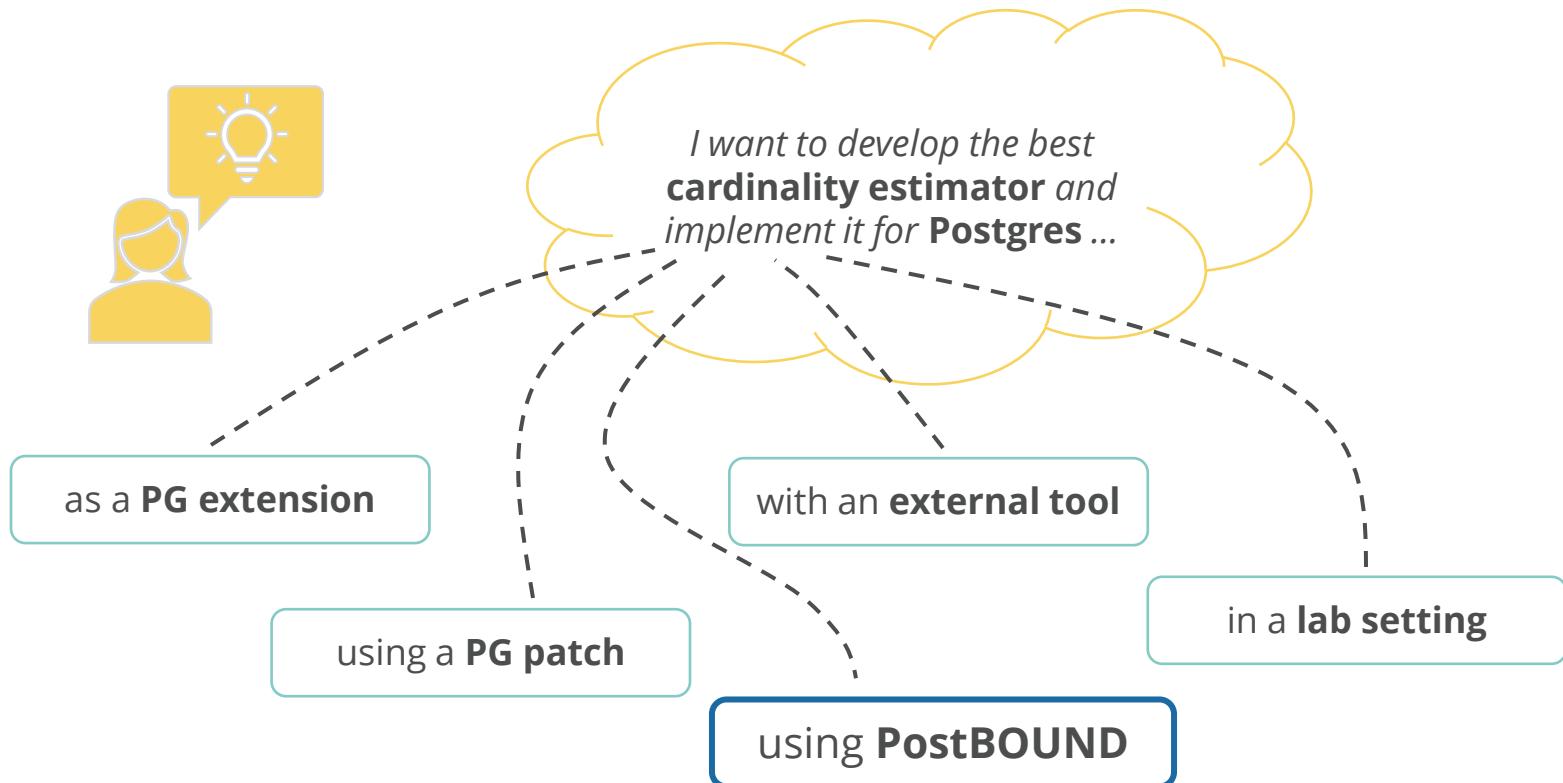
ACM 978-1-4503-7996-9/25/02 \$15.00 © 2025 ACM

<https://doi.org/10.1145/3799659>

Proc. ACM Manag. Data, Vol. 3, No. 1 (SIGMOD), Article 9. Publication date: February 2025.



From Idea To Implementation



Supervised Cardinality Estimation (*MSCN-light*)



MSCN in a nutshell

- Supervised cardinality estimation model
 - Queries are featurized according to
 - Their tables
 - The join graph
 - Filter predicates
 - Different sub-modules for each component
 - Modules are merged in final model

Our adaptation

- Feature each component with an embedding model
 - Predict cardinality with a simple gradient boosting model

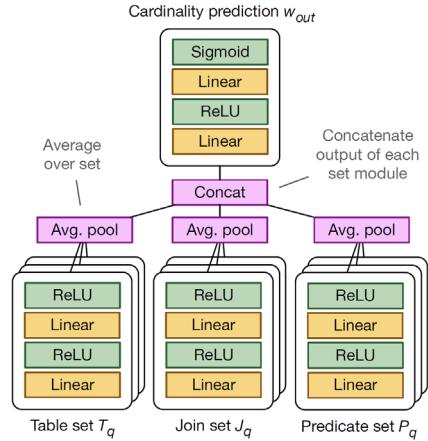
Kipf et al.: "Learned Cardinalities: Estimating Correlated Joins with Deep Learning" (CIDR'2019)

Learned Cardinalities: Estimating Correlated Joins with Deep Learning

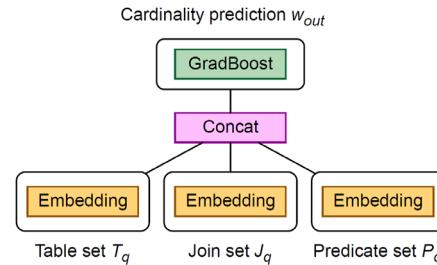
Supervised Cardinality Estimation (*MSCN-light*)



Original MSCN model



Light model



Kipf et al.: "Learned Cardinalities: Estimating Correlated Joins with Deep Learning" (CIDR'2019)



Live!

Plan Selection (*BAO-light*)



BAO in a nutshell

- Use (global) planner hints to obtain different query plans
- Vectorize query plans using binarization + schema-independent features
- Estimate plan quality using a TCNN
- Selects the (estimated) best plan for execution

Our adaptation

- Use embedding to vectorize query plan
- Predict plan cost with a simple gradient boosting model

Research Data Management Track Paper
SIGMOD '21, June 20–25, 2021, Virtual Event, China

Bao: Making Learned Query Optimization Practical

Ryan Marcus
MIT & Intel Labs
rym@mit.edu
Nesime Tatbul
MIT & Intel Labs
tatbul@csail.mit.edu
Parimkar Negi
MIT
pnegi@mit.edu
Mohammad Alizadeh
MIT
alizadeh@csail.mit.edu
Hongzhi Mao
MIT
hongzhi@mit.edu
Tim Kraska
MIT
kraska@csail.mit.edu

ABSTRACT
Recent efforts applying machine learning techniques to query optimization have shown promising results. However, these approaches often suffer from overhead, inability to adapt to changes, and poor tail performance. In this paper, we propose Bao, a learned query optimizer that can make probabilistically expensive operations in practice (thus why we wish to call it “practical”) more efficient. Bao takes advantage of the wisdom built into existing query optimizers by providing per-query optimization hints. Bao uses reinforcement learning to learn these hints. We propose a novel sampling scheme for action sampling, a well-studied reinforcement learning algorithm. As opposed to previous work, Bao does not require any prior knowledge about changes in query workload, data, and schema. Experimentally, we show that Bao outperforms traditional optimizers on several benchmarks, including end-to-end query execution performance, including tail latency for several workloads containing long-running queries. In short, Bao achieves better performance compared with a commercial system.

CCS CONCEPTS
• Information systems → Query optimization.

KEYWORDS
• Databases, machine learning, reinforcement learning

ACM Reference Format:
Ryan Marcus, Parimkar Negi, Hongzhi Mao, Nesime Tatbul, Mohammad Alizadeh, and Tim Kraska. 2021. Bao: Making Learned Query Optimization Practical. In Proceedings of the 2021 International Conference on Management of Data (SIGMOD '21), June 20–25, 2021, Virtual Event, China, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3453009.3453238>

1 INTRODUCTION
Query optimization is an important task for database management systems (DBMSs). It is a multi-step process that involves elements of query optimization – cardinality estimation and cost modeling – to produce an optimal query plan [38]. Several works have applied machine learning (ML) to query optimization [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54]. While all of these new solutions demonstrate significant improvements, they also raise the question if these techniques are yet practical, as they suffer from several fundamental problems:

- (1) Long training time. Most proposed machine learning techniques require a large amount of training data to learn, which may have a positive impact on query performance. For example, ML-powered cardinality estimators based on supervised learning require millions of training examples to learn [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54].
- (2) Inability to adjust to data and workload changes. While proposed reinforcement learning approaches can handle unpredictable changes in query workload, data, or schema can make them less effective. Reinforcement learning approaches can become less effective when data changes, or risk becoming stale. Several proposed reinforcement learning techniques assume that both the training and testing environments are identical. This assumption is violated when this is not the case [40, 31, 32, 39].
- (3) Tail latency. Reinforcement learning-based query optimization techniques can outperform traditional optimizers on average, but often perform catastrophically (e.g., 10x regression in query performance) when the environment changes significantly between training data vs. query. While some approaches offer statistical guarantees for the performance of reinforcement learning approaches, even if rare, are unacceptable in many real world applications.
- (4) Lack of explainability. Reinforcement learning approaches are already complex, understanding query optimization is even harder when black box deep learning approaches are used. More research is needed to understand how learned query optimizers do not provide a way for database administrators to influence the query optimizer’s behavior.
- (5) Integrations costs. To the best of our knowledge, all previous learned optimizers are still research prototypes, offering little to no integration with existing DBMSs. Most DBMSs support a subset of standard SQL, not to mention vendor specific features. Hence, fully integrating learned optimizers into DBMSs is a non-trivial task. The database system is not a trivial undertaking.

To the best of our knowledge, Bao (Bridging optimization in the fast lane) is the first learned query optimizer that addresses all of these problems. Bao is fully integrated into PostgreSQL, as an extension, and can be used by any DBMS administrator. The DBMS administrator (DBA) just needs to download our open-source module¹, and even has the option to selectively turn the learned optimizer on or off for specific queries.

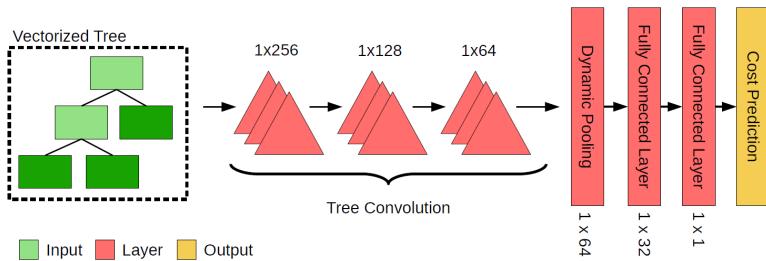
¹<https://bao.csail.mit.edu>

Marcus et al.: “BAO: Making Learned Query Optimization Practical” (SIGMOD’2021)

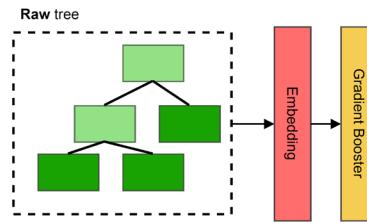
Plan Selection (*BAO-light*)



Original BAO model



Light model



Marcus et al.: "BAO: Making Learned Query Optimization Practical" (SIGMOD'2021)



Live!

Beyond Learned Optimizers: Pessimistic



Core idea

- Use upper bounds instead of cardinality point-estimates to derive a robust execution plan
- Disable dangerous join operators, e.g., nested-loop join

Our adaptation

- Use the UES bound formula
- Only enable hash joins
- DBMS is free to select a join order
- DBMS is free to select scan operators

Hertzschuch et al.: "Simplicity Done Right for Join Ordering" (CIDR'2021)

Simplicity Done Right for Join Ordering

Axel Hertzschuch, Claudio Hartmann, Dirk Habich, Wolfgang Lehner
Database Systems Group
Technische Universität Dresden
Dresden, Germany
firstname.lastname@tu-dresden.de

ABSTRACT
In this paper, we propose a simple, yet fast and effective, approach for join ordering selection in query execution. Our scheme comprises three building blocks: (i) a simple upper bound for arbitrary multi-join queries, (ii) a sketch-based approach for the cardinality bound, and (iii) sampling as query execution to provide tight bounds for the cardinality of join results. As we are going to show, using the Join-Order-Benchmark, our approach is able to outperform state-of-the-art systems with significantly less optimization overhead, resulting in a reduction of execution time for all 111 JOIN queries compared to state-of-the-art and recent approaches.

1. INTRODUCTION
Although query optimization has been a core research topic for decades, it is still far from being solved [7]. This is particularly true for join ordering, one of the most important steps in query optimization: finding a good join order [2, 7]. To date, many different approaches have been proposed. Many of them are required [2, 7]. This includes joins over intermediate relations, which are often used to answer complex queries. The question whether it is even possible to achieve such estimates without join execution is yet to be answered [7].
A common approach for join cardinality estimation is to rely on basic heuristics that may assume predicates independent and a uniform distribution of values. More sophisticated approaches for the join cardinality estimation have been proposed [8, 11, 21]. These approaches are often based on histograms appearing at first glance, for example [8, 11, 21], but they do not scale well to many joins [3, 21]. On the other hand, there are also approaches that are based on more advanced techniques [6, 20] as they model complex data characteristics. However, these approaches are often limited to filter predicate types and their training depends on executing a plethora of joins, which may take days or weeks [19].

This article is published under a Creative Commons Attribution License (CC BY). The full-text may be used and given to third parties, and reproduced as a reprint or for teaching purposes, without prior permission or charge, provided the full-text is not changed in any way, the full reference is given, and the material is not sold in any format. The full-text must not be sold in any format or provided to companies in any direct or indirect manner. All materials to reproduce and further analyse the results reported in this article can be found at: <https://gitlab.com/ahertz/SimplicityDoneRight>.

2. U-BLOCKS: JOIN ORDERING
The first building block U includes a simple upper bound for the cardinality of a join. For that, we assume independence of the join predicates (e.g., $\text{age} < \text{height}$ statistics) and precise filter selectivity estimates. While histograms may suffice for base filters, Section 4 shows that histograms are not sufficient for join filters (e.g., $\text{filter}(\text{join}(\text{t1}, \text{t2}))$).
2.1. U-Block: Simple Upper Bound for Joins
The second building block F includes a simple upper bound for the cardinality of a join. For that, we assume independence of the join predicates (e.g., $\text{age} < \text{height}$ statistics) and precise filter selectivity estimates. While histograms may suffice for base filters, Section 4 shows that histograms are not sufficient for join filters (e.g., $\text{filter}(\text{join}(\text{t1}, \text{t2}))$).



Live!

Reproducible Prototyping of Query Optimizer Components

Rico Bergmann and Dirk Habich

Technische Universität Dresden, Germany

SIGMOD'25 Tutorial, June 27, 2025, Berlin, Germany