



Hands-On Material:

github.com/db-tu-dresden/SIGMOD25-OptimizerTutorial

Reproducible Prototyping of Query Optimizer Components

Rico Bergmann and Dirk Habich

Technische Universität Dresden, Germany



SIGMOD'25 Tutorial, June 27, 2025, Berlin, Germany

Presenters

Rico Bergmann



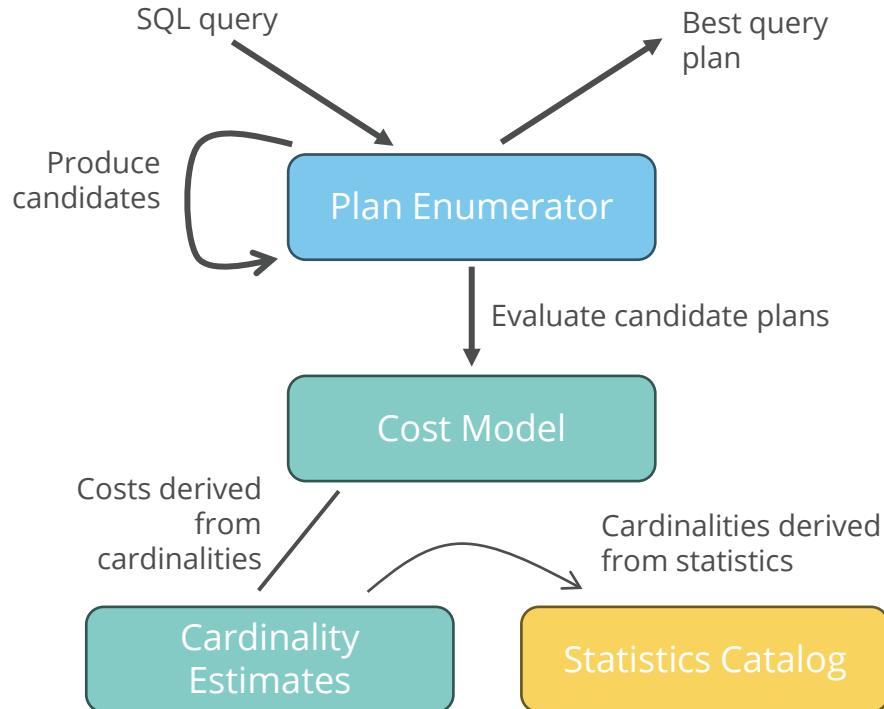
- PhD student (2nd year)
- Technische Universität Dresden, Germany
- Research focus
 - Optimizing the query optimizer
 - Developing of a new query optimizer framework

Dirk Habich



- Associate Professor
- Technische Universität Dresden, Germany
- Research focus
 - Database systems on modern hardware
 - Query optimization
 - Habilitation in 2018 on „In-Memory Database Query Processing on Large-Scale Multiprocessor Systems“

Textbook Query Optimizer



Current Research – Cost Model

Zero-Shot Cost Models for Out-of-the-box Learned Cost Prediction

Borjanan Härkönen

Teknisk Universitet af Danmark
Institut for Data Science og Data Engineering

Carter Baumig

Teknisk Universitet af Danmark / ØREF
Østasiatisk Institutt for Data Science

ABSTRACT

In this paper, we propose a novel cost model, which makes use of learned cost estimation that generalizes to unseen datasets. In contrast to previous work, our approach does not require a large number of training examples per every new dataset. Instead, it uses a small set of training examples from a single dataset to learn a general cost model that can be applied to any dataset.

To validate our approach, we sought a real-world application where a learned cost model can be used to predict the cost of training machine learning models. We chose to focus on the task of learning to predict the cost of training deep learning models, as this is a challenging task that requires a large amount of training data to build a model that performs well across many different datasets.

We find that our learned cost model can predict the cost of training a wide variety of deep learning models without requiring a large amount of training data. This is particularly important as training a learned cost model can be a time-consuming process in the machine learning pipeline.

The learned cost model can be used to fine-tune the machine learning pipeline to reduce the overall cost of training a machine learning model.

bioRxiv preprint doi: <https://doi.org/10.1101/275129>; this version posted December 1, 2017. The copyright holder for this preprint (which was not certified by peer review) is the author/funder, who has granted bioRxiv a license to display the preprint in perpetuity. It is made available under aCC-BY-NC-ND 4.0 International license.

INTRODUCTION

1.1. Database Systems

1.1.1. Database Management Systems

1.1.1.1. Database Management Systems

1.1.1.1.1. Database Management Systems

1.1.1.1.1.1. Database Management Systems

1.1.1.1.1.2. Database Management Systems

1.1.1.1.3. Database Management Systems

1.1.1.2. Database Management Systems

1.1.1.3. Database Management Systems

1.1.1.4. Database Management Systems

1.1.1.5. Database Management Systems

1.1.1.6. Database Management Systems

1.1.1.7. Database Management Systems

1.1.1.8. Database Management Systems

1.1.1.9. Database Management Systems

1.1.1.10. Database Management Systems

1.1.1.11. Database Management Systems

1.1.2. Database Systems

1.1.2.1. Database Systems

1.1.2.1.1. Database Systems

1.1.2.1.1.1. Database Systems

1.1.2.1.1.1.1. Database Systems

1.1.2.1.1.1.2. Database Systems

1.1.2.1.1.3. Database Systems

1.1.2.1.2. Database Systems

1.1.2.1.3. Database Systems

1.1.2.1.4. Database Systems

1.1.2.1.5. Database Systems

1.1.2.1.6. Database Systems

1.1.2.1.7. Database Systems

1.1.2.1.8. Database Systems

1.1.2.1.9. Database Systems

1.1.2.1.10. Database Systems

1.1.2.1.11. Database Systems

1.1.2.1.12. Database Systems

1.1.3. Database Systems

1.1.3.1. Database Systems

1.1.3.1.1. Database Systems

1.1.3.1.1.1. Database Systems

1.1.3.1.1.1.1. Database Systems

1.1.3.1.1.1.2. Database Systems

1.1.3.1.1.3. Database Systems

1.1.3.1.2. Database Systems

1.1.3.1.3. Database Systems

1.1.3.1.4. Database Systems

1.1.3.1.5. Database Systems

1.1.3.1.6. Database Systems

1.1.3.1.7. Database Systems

1.1.3.1.8. Database Systems

1.1.3.1.9. Database Systems

1.1.3.1.10. Database Systems

1.1.3.1.11. Database Systems

1.1.3.1.12. Database Systems

1.1.4. Database Systems

1.1.4.1. Database Systems

1.1.4.1.1. Database Systems

1.1.4.1.1.1. Database Systems

1.1.4.1.1.1.1. Database Systems

1.1.4.1.1.1.2. Database Systems

1.1.4.1.1.3. Database Systems

1.1.4.1.2. Database Systems

1.1.4.1.3. Database Systems

1.1.4.1.4. Database Systems

1.1.4.1.5. Database Systems

1.1.4.1.6. Database Systems

1.1.4.1.7. Database Systems

1.1.4.1.8. Database Systems

1.1.4.1.9. Database Systems

1.1.4.1.10. Database Systems

1.1.4.1.11. Database Systems

1.1.4.1.12. Database Systems

1.1.5. Database Systems

1.1.5.1. Database Systems

1.1.5.1.1. Database Systems

1.1.5.1.1.1. Database Systems

1.1.5.1.1.1.1. Database Systems

1.1.5.1.1.1.2. Database Systems

1.1.5.1.1.3. Database Systems

1.1.5.1.2. Database Systems

1.1.5.1.3. Database Systems

1.1.5.1.4. Database Systems

1.1.5.1.5. Database Systems

1.1.5.1.6. Database Systems

1.1.5.1.7. Database Systems

1.1.5.1.8. Database Systems

1.1.5.1.9. Database Systems

1.1.5.1.10. Database Systems

1.1.5.1.11. Database Systems

1.1.5.1.12. Database Systems

1.1.6. Database Systems

1.1.6.1. Database Systems

1.1.6.1.1. Database Systems

1.1.6.1.1.1. Database Systems

1.1.6.1.1.1.1. Database Systems

1.1.6.1.1.1.2. Database Systems

1.1.6.1.1.3. Database Systems

1.1.6.1.2. Database Systems

1.1.6.1.3. Database Systems

1.1.6.1.4. Database Systems

1.1.6.1.5. Database Systems

1.1.6.1.6. Database Systems

1.1.6.1.7. Database Systems

1.1.6.1.8. Database Systems

1.1.6.1.9. Database Systems

1.1.6.1.10. Database Systems

1.1.6.1.11. Database Systems

1.1.6.1.12. Database Systems

1.1.7. Database Systems

1.1.7.1. Database Systems

1.1.7.1.1. Database Systems

1.1.7.1.1.1. Database Systems

1.1.7.1.1.1.1. Database Systems

1.1.7.1.1.1.2. Database Systems

1.1.7.1.1.3. Database Systems

1.1.7.1.2. Database Systems

1.1.7.1.3. Database Systems

1.1.7.1.4. Database Systems

1.1.7.1.5. Database Systems

1.1.7.1.6. Database Systems

1.1.7.1.7. Database Systems

1.1.7.1.8. Database Systems

1.1.7.1.9. Database Systems

1.1.7.1.10. Database Systems

1.1.7.1.11. Database Systems

1.1.7.1.12. Database Systems

1.1.8. Database Systems

1.1.8.1. Database Systems

1.1.8.1.1. Database Systems

1.1.8.1.1.1. Database Systems

1.1.8.1.1.1.1. Database Systems

1.1.8.1.1.1.2. Database Systems

1.1.8.1.1.3. Database Systems

1.1.8.1.2. Database Systems

1.1.8.1.3. Database Systems

1.1.8.1.4. Database Systems

1.1.8.1.5. Database Systems

1.1.8.1.6. Database Systems

1.1.8.1.7. Database Systems

1.1.8.1.8. Database Systems

1.1.8.1.9. Database Systems

1.1.8.1.10. Database Systems

1.1.8.1.11. Database Systems

1.1.8.1.12. Database Systems

1.1.9. Database Systems

1.1.9.1. Database Systems

1.1.9.1.1. Database Systems

1.1.9.1.1.1. Database Systems

1.1.9.1.1.1.1. Database Systems

1.1.9.1.1.1.2. Database Systems

1.1.9.1.1.3. Database Systems

1.1.9.1.2. Database Systems

1.1.9.1.3. Database Systems

1.1.9.1.4. Database Systems

1.1.9.1.5. Database Systems

1.1.9.1.6. Database Systems

1.1.9.1.7. Database Systems

1.1.9.1.8. Database Systems

1.1.9.1.9. Database Systems

1.1.9.1.10. Database Systems

1.1.9.1.11. Database Systems

1.1.9.1.12. Database Systems

1.1.10. Database Systems

1.1.10.1. Database Systems

1.1.10.1.1. Database Systems

1.1.10.1.1.1. Database Systems

1.1.10.1.1.1.1. Database Systems

1.1.10.1.1.1.2. Database Systems

1.1.10.1.1.3. Database Systems

1.1.10.1.2. Database Systems

1.1.10.1.3. Database Systems

1.1.10.1.4. Database Systems

1.1.10.1.5. Database Systems

1.1.10.1.6. Database Systems

1.1.10.1.7. Database Systems

1.1.10.1.8. Database Systems

1.1.10.1.9. Database Systems

1.1.10.1.10. Database Systems

1.1.10.1.11. Database Systems

1.1.10.1.12. Database Systems

1.1.11. Database Systems

1.1.11.1. Database Systems

1.1.11.1.1. Database Systems

1.1.11.1.1.1. Database Systems

1.1.11.1.1.1.1. Database Systems

1.1.11.1.1.1.2. Database Systems

1.1.11.1.1.3. Database Systems

1.1.11.1.2. Database Systems

1.1.11.1.3. Database Systems

1.1.11.1.4. Database Systems

1.1.11.1.5. Database Systems

1.1.11.1.6. Database Systems

1.1.11.1.7. Database Systems

1.1.11.1.8. Database Systems

1.1.11.1.9. Database Systems

1.1.11.1.10. Database Systems

1.1.11.1.11. Database Systems

1.1.11.1.12. Database Systems

1.1.12. Database Systems

1.1.12.1. Database Systems

1.1.12.1.1. Database Systems

1.1.12.1.1.1. Database Systems

1.1.12.1.1.1.1. Database Systems

1.1.12.1.1.1.2. Database Systems

1.1.12.1.1.3. Database Systems

1.1.12.1.2. Database Systems

1.1.12.1.3. Database Systems

1.1.12.1.4. Database Systems

1.1.12.1.5. Database Systems

1.1.12.1.6. Database Systems

1.1.12.1.7. Database Systems

1.1.12.1.8. Database Systems

1.1.12.1.9. Database Systems

1.1.12.1.10. Database Systems

1.1.12.1.11. Database Systems

1.1.12.1.12. Database Systems

1.1.13. Database Systems

1.1.13.1. Database Systems

1.1.13.1.1. Database Systems

1.1.13.1.1.1. Database Systems

1.1.13.1.1.1.1. Database Systems

1.1.13.1.1.1.2. Database Systems

1.1.13.1.1.

LEON: A New Framework for ML-Aided Query Optimization

Xia Chen
University of Chinese
Academy of Sciences
and Technology of
China
ml-optimizer.cs.ac.cn

Jianglong Wang
Huawei Noah's Ark
Laboratory
ml-optimizer.cs.ac.cn

Hafizan Chen
University of Chinese
Academy of Sciences
and Technology of
China
hafizan@cs.ac.cn

Kai Zeng
Huawei Noah's Ark
Laboratory
ml-optimizer.cs.ac.cn

Zhuo Liang
University of Chinese
Academy of Sciences
and Technology of
China
zhuoliang@cs.ac.cn

Han Su
University of Chinese
Academy of Sciences
and Technology of
China
han.su@cs.ac.cn

Shengming Liu
University of Chinese
Academy of Sciences
and Technology of
China
shengming.liu@cs.ac.cn

Kai Zheng
University of Chinese
Academy of Sciences
and Technology of
China
kai.zheng@cs.ac.cn

ABSTRACT

Query optimization has long been a fundamental and challenging problem in database systems. In this paper, we propose a new framework, named LEON, for query optimization. LEON integrates machine learning (ML) into query optimization, and it can automatically learn the ML models from historical execution traces to predict the performance of different query plans. LEON can also automatically generate the best query plan for a given query by selecting the most promising one from the predicted results.

INTRODUCTION

The query optimizer is a core component of modern database management systems (DBMS). The main purpose of the query optimizer is to find the optimal query execution plan for a given query. The query optimizer needs to consider many factors, such as the query complexity [3], the query statistics [1], the system configuration [2], and the user requirements [4].

Traditionally, the query optimizer is implemented by hand-coded rules and heuristics. Although these methods have been effective, they are difficult to maintain and update. In recent years, there has been a growing interest in using machine learning (ML) to assist query optimization. ML-based query optimizers have been proposed to predict the performance of different query plans [5, 6]. These approaches have shown promising results, but they still have some limitations. For example, they often require a large amount of training data, which is difficult to obtain in practice. Moreover, they may not always be able to handle complex queries or handle changes in the system environment.

In this paper, we propose a new framework, named LEON, for query optimization. LEON integrates ML into query optimization, and it can automatically learn the ML models from historical execution traces to predict the performance of different query plans.

The rest of the paper is organized as follows. Section 2 gives an overview of related work. Section 3 presents the details of LEON. Section 4 evaluates the performance of LEON. Section 5 concludes the paper.

Neo: A Learned Query Optimizer

Ryan Marcus¹, Parhamineh Negi², Hongbo Mao², Chi Zhang²,
Hamidreza Alizadeh³, Tim Kraska⁴, Ciprian Popescu⁵, Nesrine Tafti⁶
¹ml-optimizer.cs.ac.cn, ²ml-optimizer.cs.ac.cn,
³ml-optimizer.cs.ac.cn, ⁴ml-optimizer.cs.ac.cn,
⁵ml-optimizer.cs.ac.cn, ⁶ml-optimizer.cs.ac.cn

Abstract. Neo is one of the most challenging problems in the DeepMind competition over the last decades. In this paper, we propose a learned query optimizer, named Neo, which can automatically learn the query optimizer and predict the query performance. It is the first learned query optimizer that can handle complex queries and can be applied to real-world scenarios. Neo can automatically generate the best query plan for a given query by selecting the most promising one from the predicted results.

We evaluate the performance of Neo on a real-world dataset and compare it with other learned query optimizers. The results show that Neo can achieve better performance than other learned query optimizers. We also compare Neo with a hand-coded query optimizer and show that Neo can achieve similar performance to the hand-coded query optimizer.

Keywords: Learned query optimizer, DeepMind competition, Machine learning, Query optimization, Database management system.

ABSTRACT

A learned query optimizer is one of the most challenging problems in the DeepMind competition over the last decades. In this paper, we propose a learned query optimizer, named Neo, which can automatically learn the query optimizer and predict the query performance. It is the first learned query optimizer that can handle complex queries and can be applied to real-world scenarios. Neo can automatically generate the best query plan for a given query by selecting the most promising one from the predicted results.

We evaluate the performance of Neo on a real-world dataset and compare it with other learned query optimizers. The results show that Neo can achieve better performance than other learned query optimizers. We also compare Neo with a hand-coded query optimizer and show that Neo can achieve similar performance to the hand-coded query optimizer.

Keywords: Learned query optimizer, DeepMind competition, Machine learning, Query optimization, Database management system.

PCMR Reference Format:
Myqiao Zhou: Neo: A Learned Query Optimizer. *Proc. VLDB Endow.* 14, 12 (2021), 147–160.

1 Equal contribution. Corresponding author.

2 Department of Computer Science, University of Chinese Academy of Sciences, Beijing, China.
3 Department of Computer Science, Tsinghua University, Beijing, China.
4 Department of Computer Science, Stanford University, Stanford, CA, USA.
5 Department of Computer Science, University of Illinois Urbana-Champaign, IL, USA.
6 Department of Computer Science, University of Michigan, Ann Arbor, MI, USA.

✉ myqiao.zhou@cs.ac.cn (M. Zhou); hongbo.mao@cs.ac.cn (H. Mao); chi.zhang@cs.ac.cn (C. Zhang); nesrine.tafti@umich.edu (N. Tafti).

✉ Correspondence to: Ciprian Popescu (ciprian.popescu@uiuc.edu).

Received: 2020-01-01; revised: 2020-06-01; accepted: 2020-07-01.

© The Author(s) 2021. Published by the VLDB Endowment.

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License (CC BY-NC-SA).

For more information, see <http://creativecommons.org/licenses/by-nc-sa/4.0/>.

DOI: <https://doi.org/10.14778/3459902.3460000>

Published online in VLDB Endowment at www.vldb.org/pv2021.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

VLDB Endowment is a registered trademark of the VLDB Endowment.

<

Current Research



Join Ordering

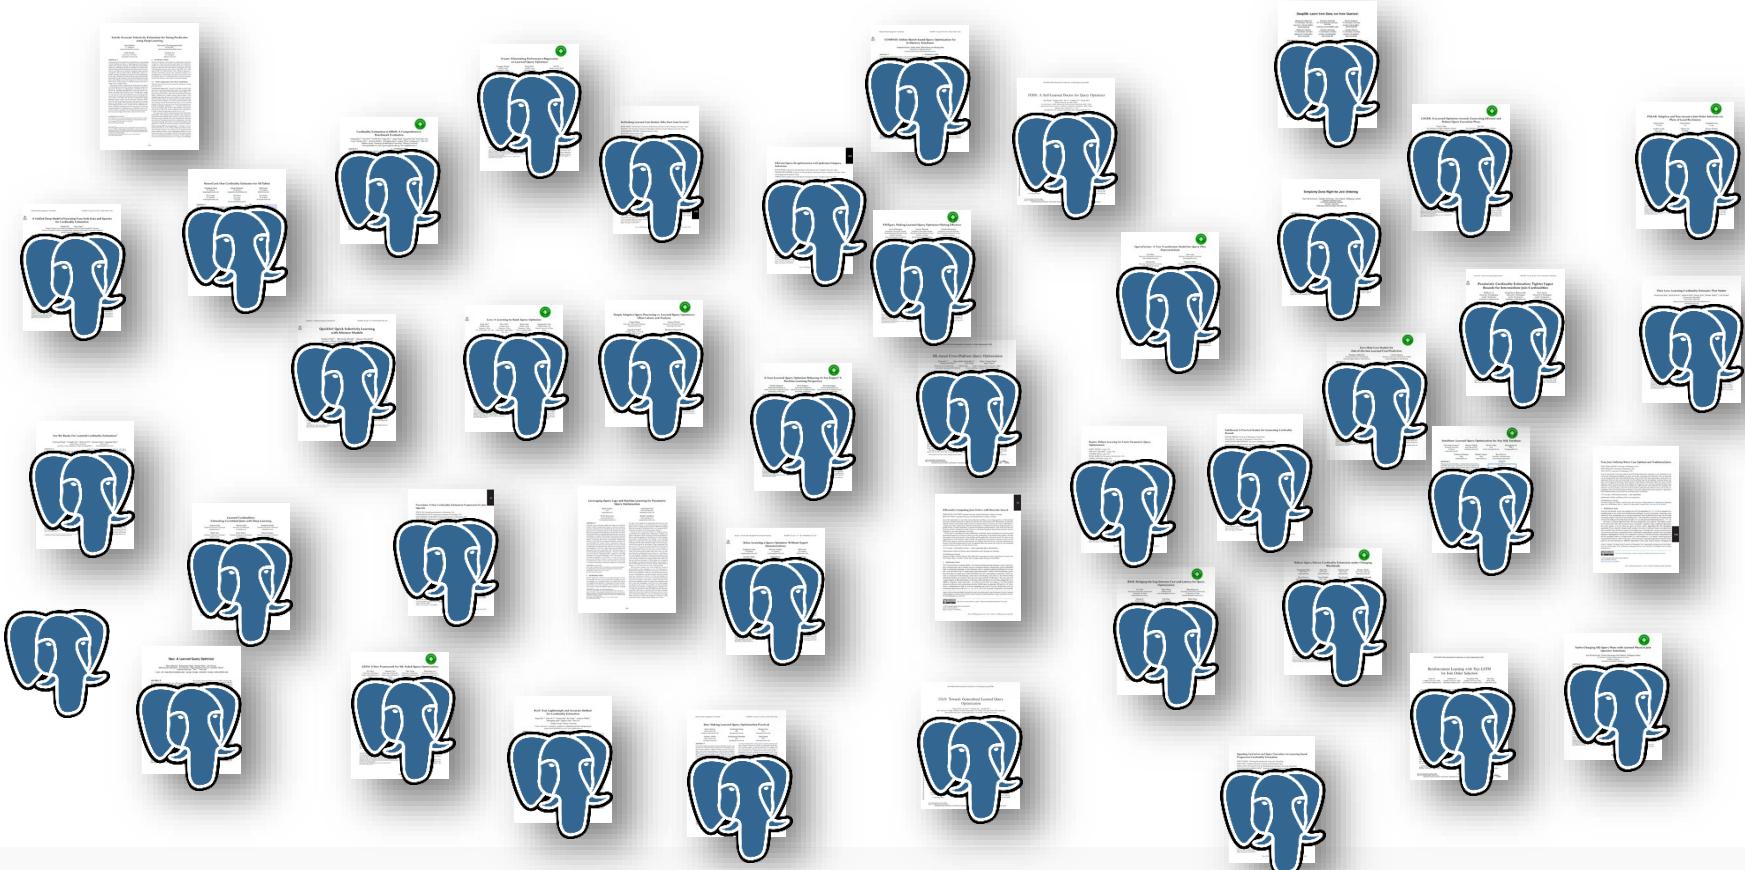
Operator Selection

Pessimistic Optimization

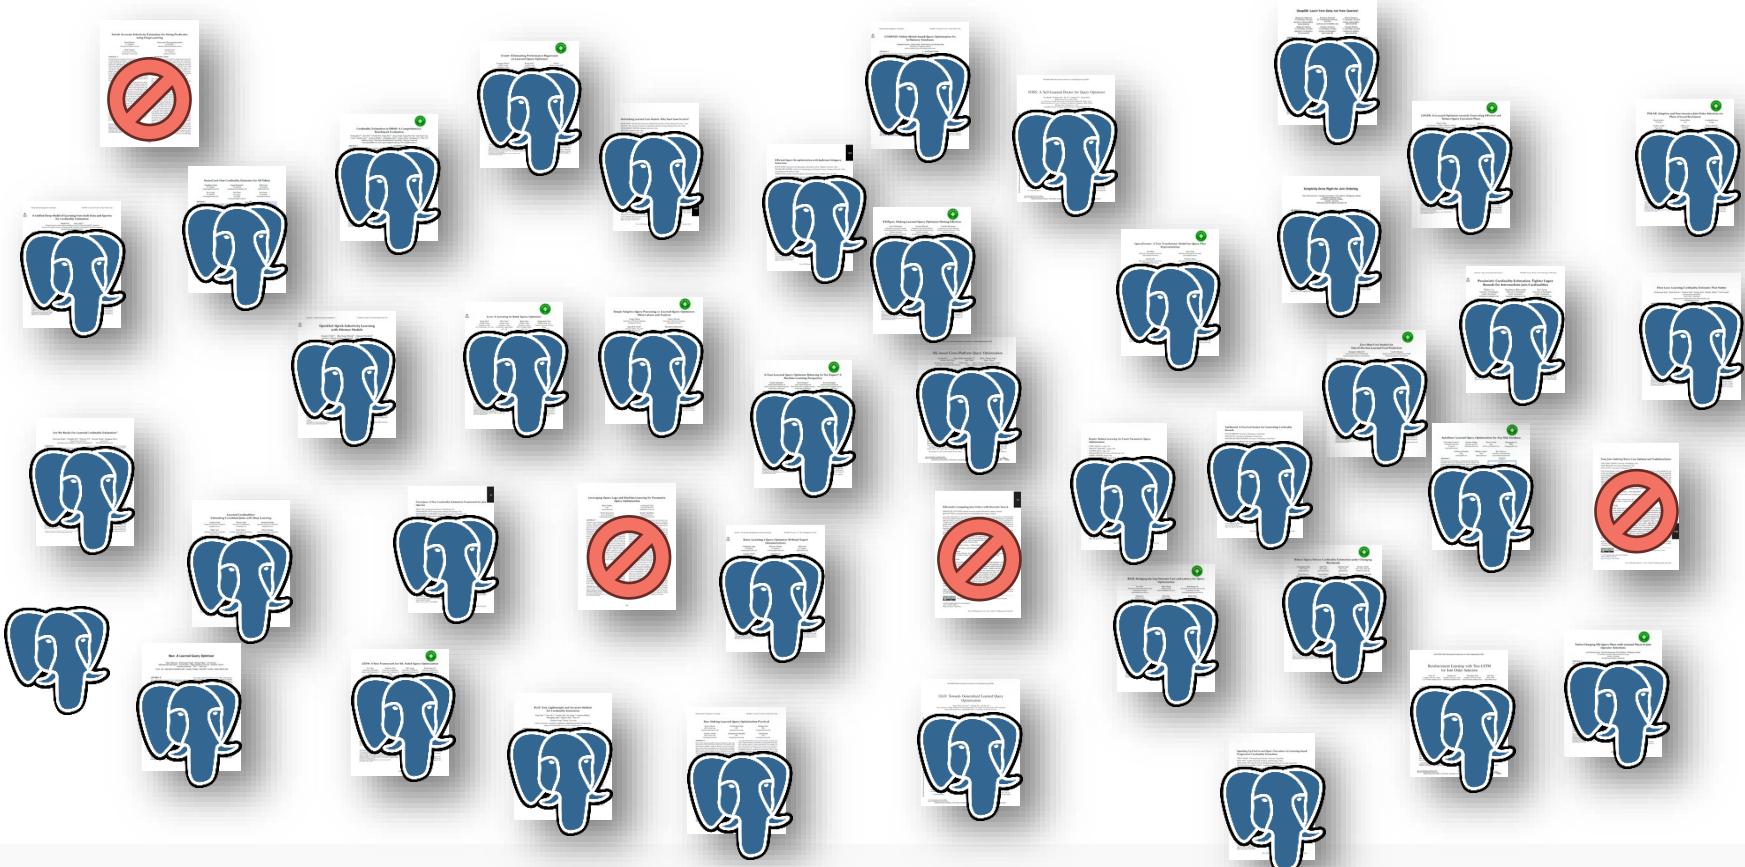
Current Research – Common Ground



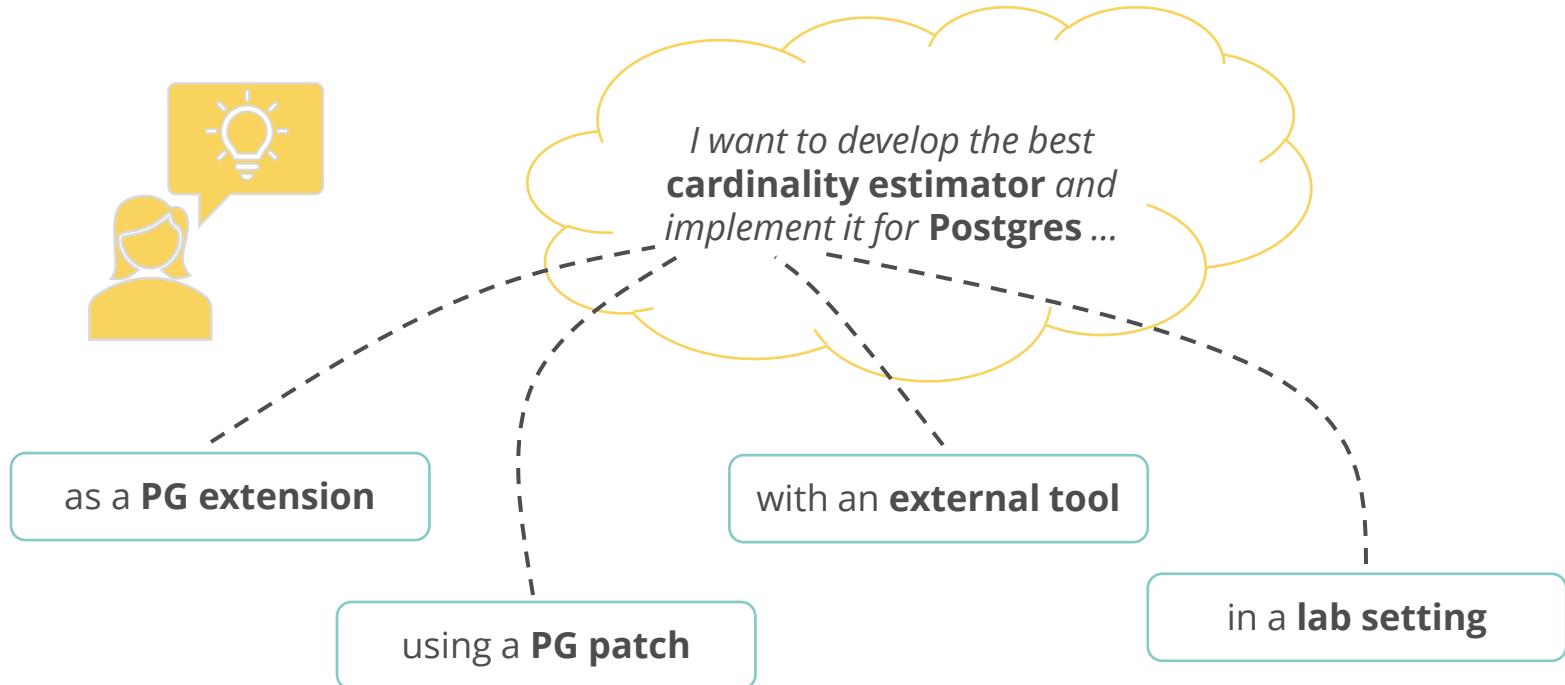
Current Research – Common Ground



Current Research – Common Ground



From Idea To Implementation



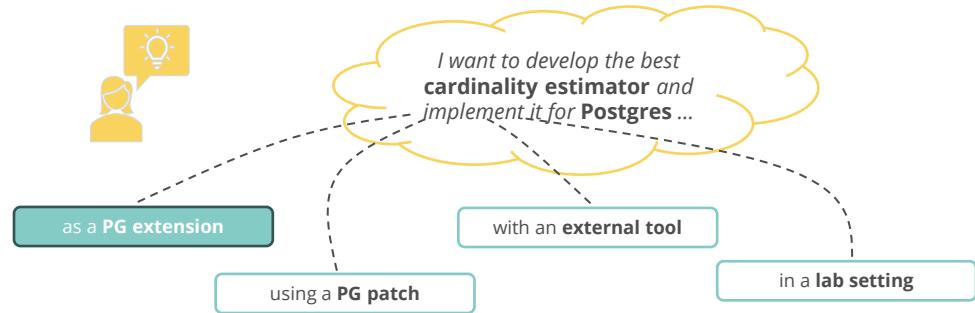
Implementation

Postgres Extension

- Postgres includes powerful extension mechanism
- Dynamically added to the server process
- Provides abstract functions to plug in custom functionality

Existing Hooks

- Complete optimization
 - planner_hook
- Join ordering + operator selection
 - join_search_hook
- Operator selection
 - set_rel_pathlist_hook
 - set_join_pathlist_hook
 - create_upper_paths_hook



Downsides

- Relies on available hooks
 - E.g., currently no hook for cardinality estimator or cost model
- Requires deep understanding of PG internals

```
double max_column_frequency(PlannerInfo *root, RelOptInfo *rel, Var *column) {
    AttStatsSlot sslot;
    VariableStatData vardata;
    double max_freq;
    examine_variable(root, (Node *) column, 0, &vardata);
    get_attstatsslot(&sslot, vardata.statsTuple, STATISTIC_KIND_MCV, InvalidOid,
                      ATTSTATSSLOT_VALUES | ATTSTATSSLOT_NUMBERS);

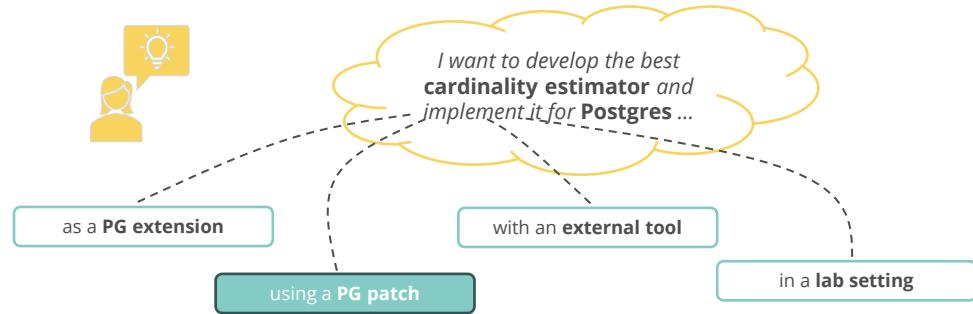
    if (sslot.nvalues > 0)
        /* We found an MCV list, use it */
        max_freq = sslot.numbers[0] * rel->tuples; /* index 0 is the highest frequency */
    else {
        /* No MCV found, assume equal distribution */
        bool default_est; /* ignored, keep compiler quiet */
        double ndistinct = get_variable_numdistinct(&vardata, &default_est);
        if (ndistinct < 0) /* negative means ND is a fraction of the table size */
            ndistinct *= -1.0 * rel->tuples;
        max_freq = rel->tuples / ndistinct;
    }

    ReleaseVariableStats(vardata);
    return max_freq;
}
```

Implementation

Postgres Patch

- Direct modifications of the PG source code
 - Enables changes to query optimizer and execution engine
- Adaptive query optimization becomes possible



```
8347  /*  
8348   * clamp_row_est  
8349 @@ -4733,6 +4758,10 @@ set_joinrel_size_estimates(PlannerInfo *root, RelOptInfo *rel,  
8350  
8351  
8352  
8353 +     if (enable_lero) {  
8354 +         lero_pgsql_set_joinrel_size_estimates(root, rel, outer_rel,  
8355 +                                         inner_rel, sjinfo, restrictlist);  
8356 +     }  
8357 }  
8358 /*  
8359 */
```

Downsides

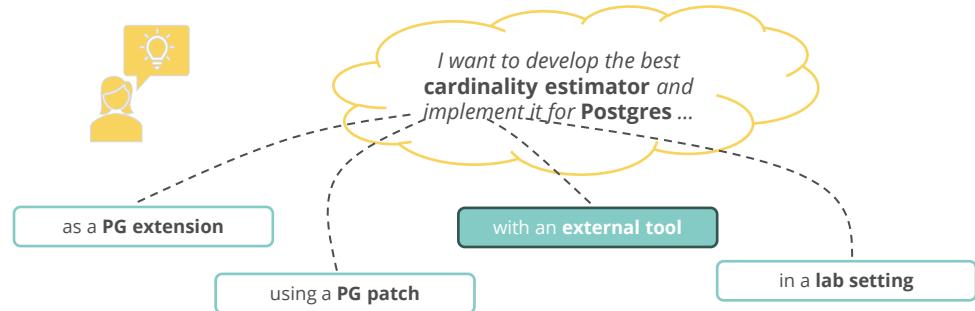
- Requires deep understanding of PG internals
- Changes are hard to track
- Specific to (minor) releases of PG

Implementation

pg_hint_plan

- Extension to embed optimizer decisions in comments
- Hints can be generated through arbitrary software
- Allows good tooling integration (e.g., ML frameworks)

```
/*+ NestLoop(t mi) IndexScan(mi) */
SELECT count(*)
FROM title t, movie_info mi
WHERE t.id = mi.movie_id
    AND t.production_year > 2010
```



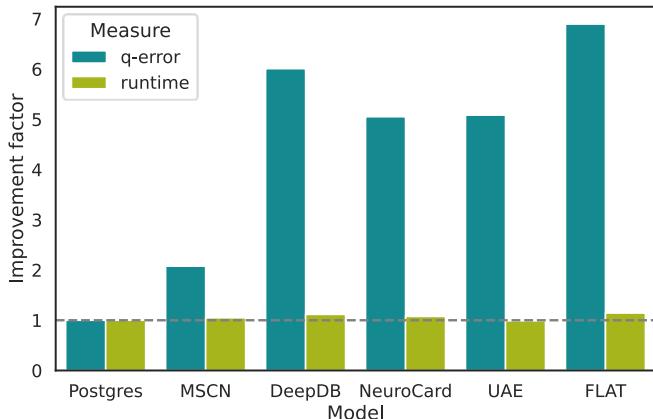
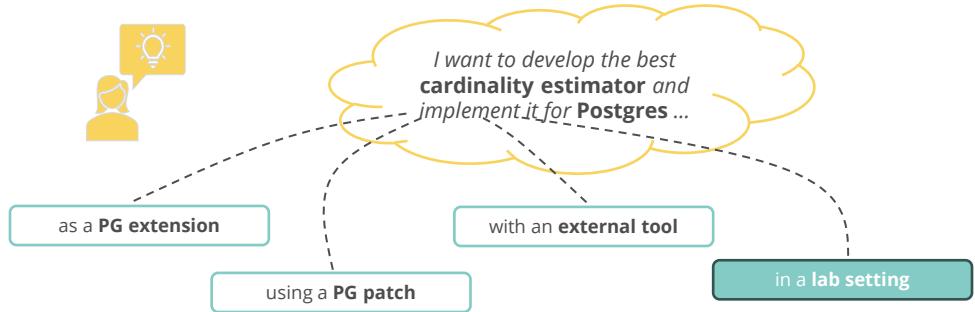
Downsides

- Opaque implementation – no standardized generation of hints
- Limited to the available hints
 - No cardinalities for base tables
 - No parallel workers

Implementation

Artificial Setting

- Don't implement for a real database
- Instead, use a "laboratory setting"
- E.g., measure q-error instead of runtime

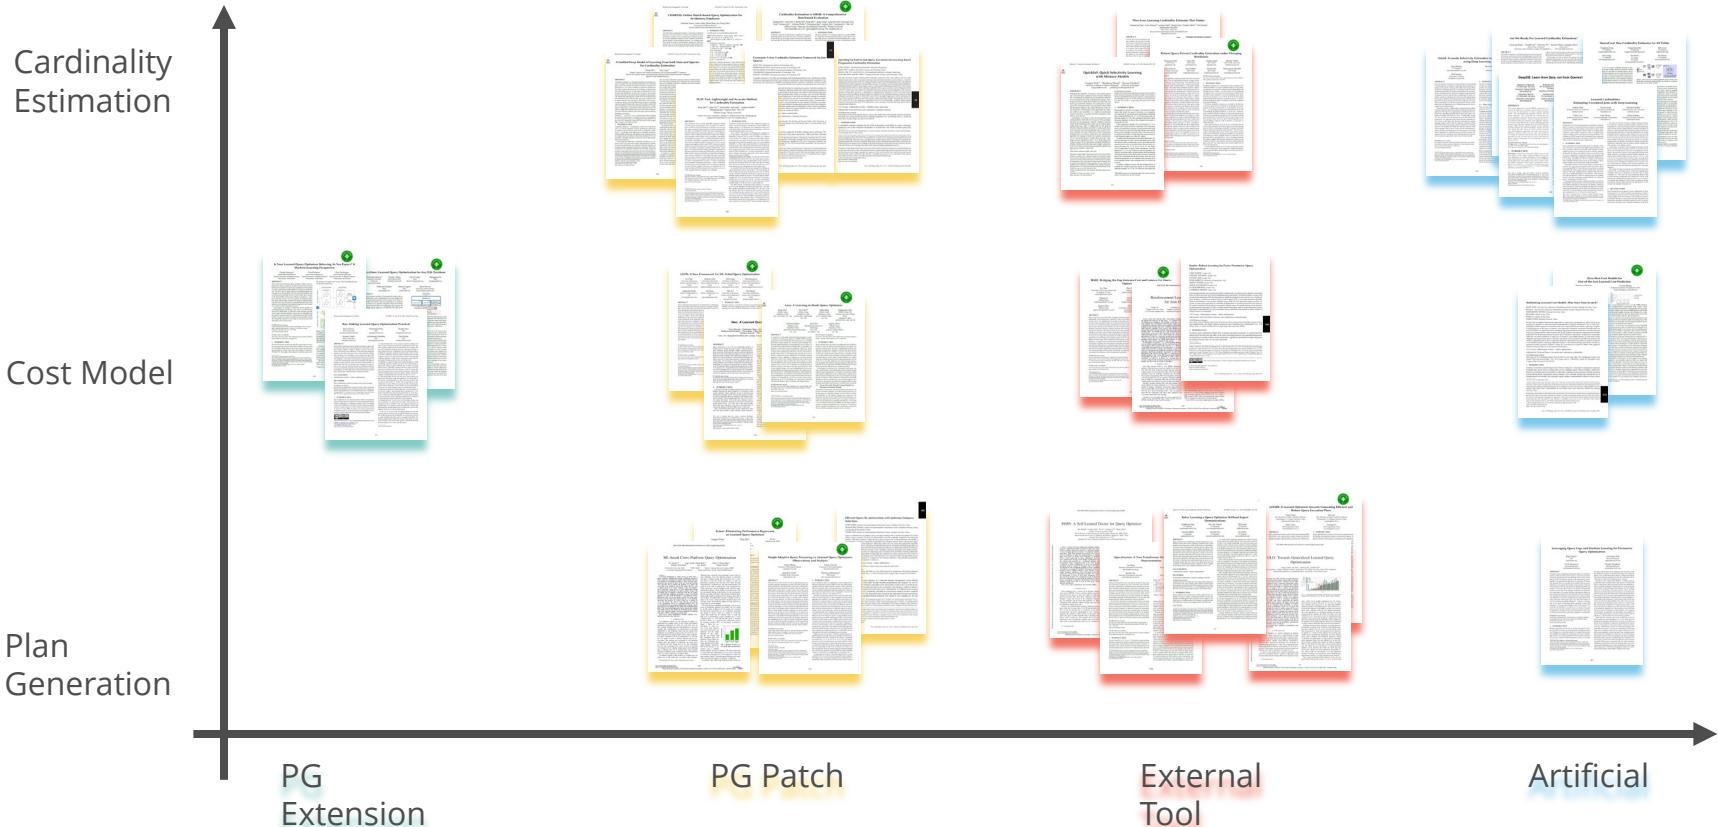


Downsides

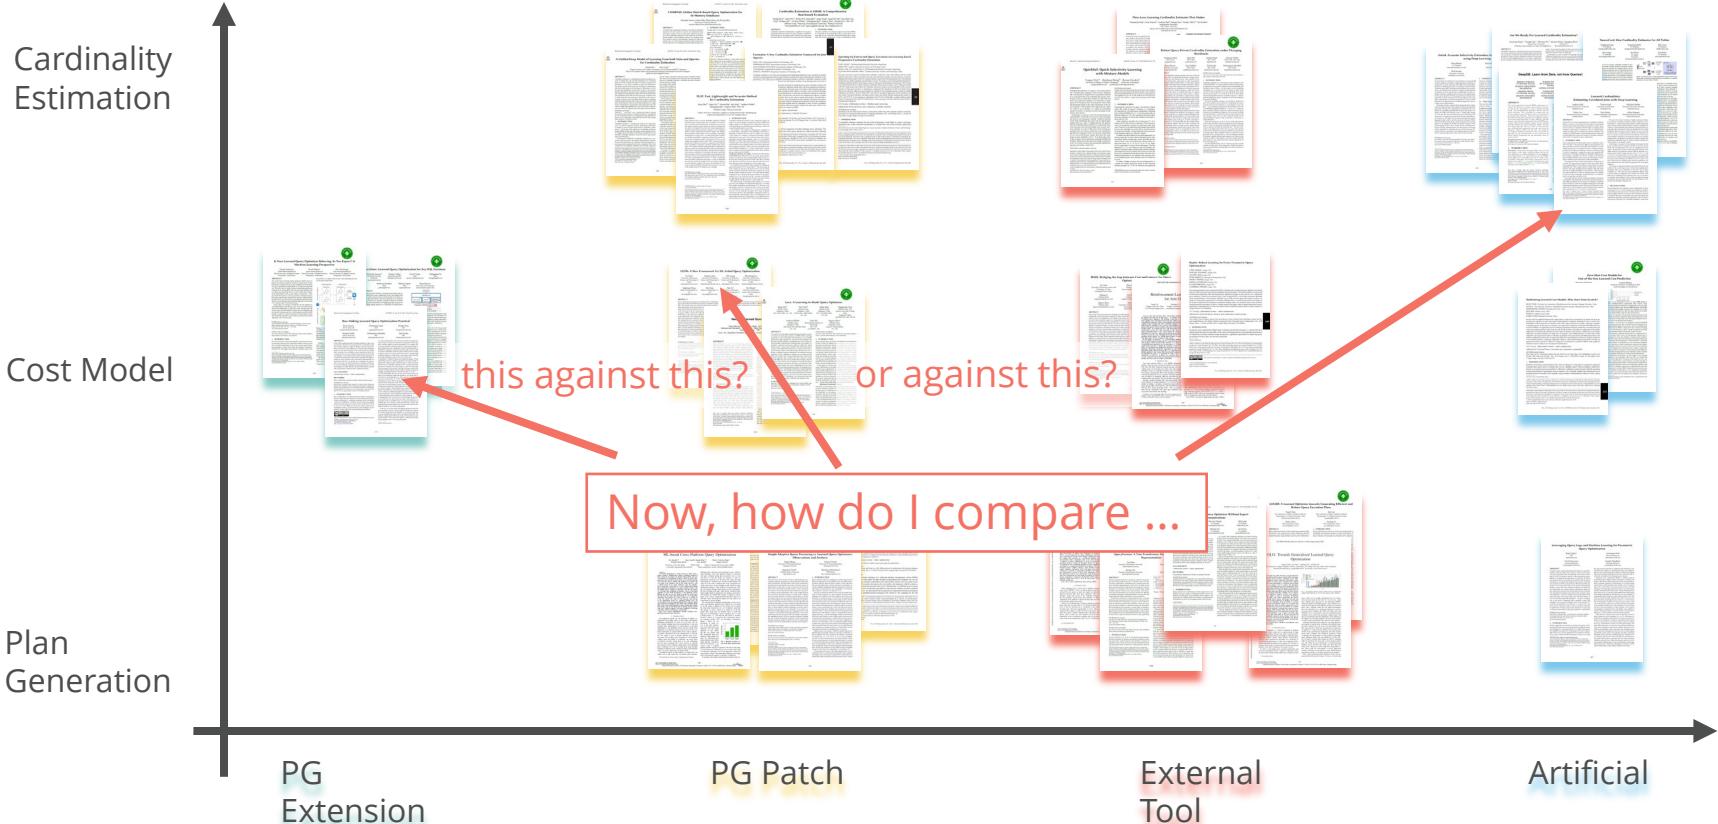
- No real-world execution – how to ensure correlation?

Numbers based on Han et al.:
Cardinality Estimation in DBMS: A Comprehensive Benchmark Evaluation (VLDB'21)

Current Research Revisited



Current Research Revisited



Let's compare with BAO (Marcus et al.)

BAO in a nutshell

- Uses search space restrictions to obtain different query plans
- Estimates the plan quality via a learned model
- Selects the (estimated) best plan for execution
- SIGMOD best paper award 2021

Implementation details

- Implemented as a PG extension
- Communicates with a Python server
- Based on PG v12.4

Marcus et al.: "BAO: Making Learned Query Optimization Practical" (SIGMOD'2021)

SIGMOD '21, June 20–25, 2021, Virtual Event, China

Research Data Management Track Paper

Bao: Making Learned Query Optimization Practical

Ryan Marcus
MIT & Intel Labs
rymanmarcus@mit.edu
Parimkar Negi
MIT
pnegi@mit.edu
Hongzhi Mao
MIT
hongzhi@mit.edu

Nesime Tatbul
MIT & Intel Labs
tatbul@csail.mit.edu
Mohammad Alizadeh
MIT
alizadeh@csail.mit.edu
Tim Kraska
MIT
kraska@csail.mit.edu

ABSTRACT
Recent efforts applying machine learning techniques to query optimization have shown promising results. However, these approaches often suffer from overhead, inability to adapt to changes, and poor tail performance. In this paper, we propose Bao, a learned query optimizer that can make probabilistically expensive operation in practice (thus why we wish to call it "learned") practical. Bao takes advantage of the window built-in existing query optimizers by providing per-query optimization hints. Bao uses a novel sampling, a well-studied reinforcement learning algorithm. As opposed to previous work, Bao does not require any prior knowledge about the system, sampling, or workload. Bao is able to handle schema changes in query workloads, data, and schemas. Experimentally, we show that Bao outperforms state-of-the-art query optimizers in terms of end-to-end query execution performance, including tail latency for several workloads containing long-running queries. In short, Bao is able to provide better performance compared with a commercial system.

CCS CONCEPTS
Information systems → Query optimization.

KEYWORDS
Database management, machine learning, reinforcement learning

ACM Reference Format:
Ryan Marcus, Parimkar Negi, Hongzhi Mao, Nesime Tatbul, Mohammad Alizadeh, and Tim Kraska. 2021. Bao: Making Learned Query Optimization Practical. In Proceedings of the 2021 International Conference on Management of Data (SIGMOD '21), June 20–25, 2021, Virtual Event, China, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3453009.3453038>

1 INTRODUCTION
Query optimization is an important task for database management systems (DBMS). It is a complex problem that requires elements of query optimization – cardinality estimation and cost modeling – to work correctly [38]. Several works have applied machine learning (ML) to query optimization [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54]. While all of these new solutions demonstrate some promise, they also raise the question that many of the techniques are yet practical, as they suffer from several fundamental problems:

- (1) Long training time. Most proposed machine learning techniques require a large amount of training data to learn, which may have a positive impact on query performance. For example, ML-powered cardinality estimators based on supervised learning require millions of training examples to learn [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54].
- (2) Inability to adjust to data and workload changes. While proposed reinforcement learning approaches are able to handle impractical changes in query workload, data, or schema can make them less useful in real-world scenarios. For example, Bao must be retrained when data changes, or risk becoming stale.
- (3) Interactions with DBMS. Reinforcement learning techniques sometimes that both the DBMS and the learned optimizer must be trained together, complete retraining when this is not the case [40, 31, 32, 39].
- (4) Tail latency. Reinforcement learning approaches often require training data to learn. While some approaches offer statistical guarantees on the quality of the learned plans, others do not. This is problematic, as even if one is unacceptable in many real world applications.
- (5) Understanding. Reinforcement learning approaches are already complex, understanding query optimization is even harder when black box deep learning approaches are used. More research is needed to understand how learned query optimizers do provide a way for database administrators to influence query execution.
- (6) Integrations costs. To the best of our knowledge, all previous learned optimizers are still research prototypes, offering little to no integration with existing DBMSs. This is a significant limitation of standard SQL, not to mention vendor specific features. Hence, fully integrating learned optimizers into a DBMS is a non-trivial task. The database system is not a trivial undertaking.

To the best of our knowledge, Bao (full name) is the first learned query optimizer that addresses all of these mentioned problems. Bao is fully integrated into PostgreSQL, as an extension, and can be used by any PostgreSQL user. The database administrator (DBA) just needs to download our open-source module¹, and even has the option to selectively turn the learned optimizer on or off for specific queries.

1 https://github.com/marcusryman/bao

This work is licensed under a Creative Commons Attribution International 4.0 License
© 2021. Copyright held by the owner/author(s).
<https://doi.org/10.1145/3453009.3453038>

1279

Let's compare with BAO (Marcus et al.)

```
~/pg-17/contrib/bao/pg_extension [master] $ make && make install
gcc -Wall -Wmissing-prototypes -Wpointer-arith -Wdeclaration-after-statement -Werror=vla
-Wendif-labels -Wmissing-format-attribute -Wimplicit-fallthrough=3 -Wcast-function-type
-Wshadow=compatible-local -Wformat-security -fno-strict-aliasing -fwrapv -fexcess-
precision=standard -Wno-format-truncation -Wno-stringop-truncation -O2 -fPIC -
fvisibility=hidden -I. -I./ -I/home/rico/pg-17/build/include/postgresql/server -
I/home/rico/pg-17/build/include/postgresql/internal      -D_GNU_SOURCE    -c -o main.o
main.c
In file included from main.c:9:
bao_bufferstate.h:13:10: fatal error: utils/reffilenodemap.h: No such file or directory
 13 | #include "utils/reffilenodemap.h"
   | ^~~~~~
compilation terminated.
make: *** [<>builtin>: main.o] Error 1
```

We can't reproduce on PG 17!

OK, what about Lero? (Zhu et al.)

Lero in a nutshell

- Uses modified cardinality estimates to obtain different query plans
- Ranks the plans using a BAO-style model
- Selects the (ranked) best plan for execution

Implementation details

- Implemented as a PG patch
- Communicates with a Python server
- Based on PG v13.1

The image shows the front page of a research paper titled "Lero: A Learning-to-Rank Query Optimizer". The page includes the title, authors, their affiliations, abstract, introduction, and some initial sections of the paper. The authors listed are Rong Zhu, Wei Chen, Bolin Ding, Xingguang Chen, Andreas Pfeller, Zirui Wu, and Jigeng Zhou. The paper is published in the Proceedings of the VLDB Endowment, Vol. 16, No. 5 ISSN 2150-8097, DOI 10.14778/3570024.3583007.

ABSTRACT
A recent trend of work apply machine learning techniques to assist query optimizers in DBMS. While exhibiting superiority in some benchmarks, their deficiencies, e.g., unreliable cardinality estimation, limit their applications. In this paper, we propose Lero, a learning-to-rank query optimizer that can learn from the inherent hardness of predicting the cost or latency of execution plans using machine learning models. In this paper, we build on top of a native query optimizer and continuous learning framework. Lero can learn from the historical execution logs to predict that the relative order or rank of plans, rather than the exact cost or latency, is sufficient for query optimization. Lero employs a pairwise learning-to-rank model to learn the relative order of two plans to predict which one is better. Such a binary classification task is much easier than a regression task. We evaluate Lero on two benchmarks. The results show that Lero can significantly outperform the native query optimizer. With its non-intrusive design, Lero can be implemented on top of any existing query optimizer. We implement Lero on PostgreSQL and demonstrate its outstanding performance using PostgreSQL. The experimental results show that Lero can reduce the average execution time by up to 37%, and other learning query optimizers can also benefit from Lero. Lero can also automatically adapt to query workloads and changes in data.

INTRODUCTION
Query optimizer plays one of the most significant roles in databases. It aims to select an efficient execution plan for each query. In general, there are two main types of query optimizers: traditional one-based query optimizers [46] and the ones with the machine learning approach [1]. Traditional query optimizers usually or other user-specified metrics about resource consumption. Such cost models contain various formulas to approximate the actual cost of executing a query. These cost models are usually manually and extensively tuned based on engineering practice. Recent work has shown that machine learning can also be used to improve query optimization algorithms with machine learning techniques. Although some progress has been made, they still suffer from deficiencies caused by the intrinsic nature of machine learning models.

In this paper, we propose Lero, a learning-to-rank query optimizer that can learn from historical execution logs to assist a query planning model for query optimization. Lero adopts a non-intrusive design, which means it can be easily integrated into any existing DBMS. Instead of building from scratch, Lero is designed to leverage decades of wisdom of query optimizers without extra potentially costly modifications. Lero can also benefit from other learning query optimizers by judiciously exploring different optimization strategies. This paper makes the following contributions:

1.1 From Heuristic-Based Costing Models to Machine Learning Models
Machine learning models have three major components: cardinality estimator, cost model, and plan enumerator. For an input query Q , cardinality estimator can be invoked to estimate the number of tuples in the result of Q . Cost model can be used to calculate the cost of a plan P with physical operators, e.g., Project , Join , Scan , etc. Plan enumerator can be used to generate other user-specified metrics regarding the efficiency of executing P . The cost model $\text{PlanCost}(P)$ estimates P 's cost and usually a function of cardinality of Q and P . The plan enumerator considers valid plans of Q in its search space and returns the one with the minimum estimated cost for execution.

Zhu et al.: "Lero: A Learning-to-Rank Query Optimizer" (VLDB'2023)

OK, what about Lero? (Zhu et al.)

```
~/pg-17 [master] $ git apply contrib/lero/0001-init-lero.patch
error: patch failed: src/backend/Makefile:19
error: src/backend/Makefile: patch does not apply
error: patch failed: src/backend/optimizer/path/costsize.c:107
error: src/backend/optimizer/path/costsize.c: patch does not apply
error: patch failed: src/backend/optimizer/plan/planner.c:64
error: src/backend/optimizer/plan/planner.c: patch does not apply
error: patch failed: src/backend/utils/misc/guc.c:98
error: src/backend/utils/misc/guc.c: patch does not apply
error: patch failed: src/include/optimizer/cost.h:67
error: src/include/optimizer/cost.h: patch does not apply
```

We can't reproduce on PG 17!

Can we at least use MSCN? (Kipf et al.)



Artificial evaluation

=

No implementation on database system

=

We can't reproduce on PG 17!

Kipf et al.: "Learned Cardinalities: Estimating Correlated Joins with Deep Learning" (CIDR'2019)

Learned Cardinalities: Estimating Correlated Joins with Deep Learning

Andreas Kipf
Technical University of Munich
kipf@in.tum.de

Thomas Kipf
University of Amsterdam
takipf@uva.nl

Bernhard Radke
Technical University of Munich
radke@in.tum.de

Viktor Lek
Technische Universität München
lekg@in.tum.de

Peter Boncz
Centrum Wiskunde & Informatica
boncz@cwi.nl

Alfred Kemper
Technische Universität München
kemper@in.tum.de

ABSTRACT

We argue that machine learning is a highly promising technique for solving the cardinality estimation problem. Learning can be used to estimate the cardinality of correlated joins by learning query features and the output being the estimated cardinality. In contrast to sampling-based cardinality estimation techniques, which are often limited to simple join types like index structures [15] and join ordering [23], the current techniques based on basic per-table statistics are not very good for large joins. However, we show that our model does not have to be perfect, it just needs to be better than the current, inaccurate sampling-based approaches. We show that our machine learning model can directly be leveraged by existing sophisticated query optimizers without requiring any other changes to the database system.

INTRODUCTION
Query optimization is fundamentally based on cardinality estimation. To be able to choose between different plan alternatives, the optimizer needs to estimate the cost of each plan, which includes the estimated result sizes. It is well known, however, that the estimates produced by all widely-used database systems are routinely wrong [1, 10, 11]. This is particularly problematic for joins, especially for foreign key joins, which are the most common type of join in the Movie Database (IMDb). French actors are more likely to participate in romantic movies than actors of other nationalities, which makes this an open area of research. One state-of-the-art proposal in this area is Index-based cardinality estimation [10]. The main idea is to build a sampling-based query qualifying model that samples against existing index structures. The problem is that this approach is not robust enough, because there are no qualifying samples to start with (i.e., under selective queries there are no table samples) or when no suitable indexes are available. Another approach is to use sampling-based cardinality estimation [24]—counting large estimation errors. A third approach is to use the proposed adoption of machine learning (ML) and specifically neural networks (deep learning), in many different applications and systems. The database community has been slow to embrace ML for cardinality estimation, but can be leveraged within data management systems. Recent research therefore investigates the use of ML for cardinality estimation [13, 20, 21, 27] and query optimization [13, 23, 27], and even indexing [12].

This article is published under a Creative Commons Attribution License. The full-text may be used and given to others, and given to appropriate repositories, without prior permission or knowledge of the publisher, provided that the full-text is not changed in any way and that the full-text gives full details of the original publication in the journal in which the article originally appeared. The full-text must not be sold in any format or supplied via the internet without the permission of the copyright holders.

We evaluate our approach using the real-world IMDB dataset [15] and show that our technique is more robust than sampling-based techniques and even is competitive in the sweet spot of these techniques, i.e., when the cardinality is between 10 and 100. This is achieved by leveraging the footprint size of about 10MB (whereas the sampling-based techniques have access to index structures and can thus sample much more effectively).

Deep learning is also well suited to query optimization by these means [13, 23, 27] that formulate join ordering as a reinforcement learning problem and use ML to find query plans. This work, in contrast, focuses on cardinality estimation, which is a much simpler task in isolation. This focus is motivated by the fact that modern join estimation algorithms can find the optimal join order for queries with dozens of relations [26]. Cardinality estimation, on the other

Could frameworks help us?

Ideally, we would like:



Broad support for different optimization approaches



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

Existing Frameworks

Research Frameworks

Apache Calcite

- Java-based framework to implement entire optimizer
- Community-driven development
- Deep integration in Java/Apache ecosystem
- Allows customization of a traditional optimizer model

<https://calcite.apache.org/>



	Broad support for different optimization approaches
	Transparent implementation of prototypes
	Execution on Postgres
	Easy comparison of different prototypes

Research Frameworks

LingoDB

<https://www.lingo-db.com/>

- Research system focused on MLIR-based query processing
- Developed at TUM
- Optimizer implemented as LLVM compiler passes



	Broad support for different optimization approaches
	Transparent implementation of prototypes
	Execution on Postgres
	Easy comparison of different prototypes

Research Frameworks

PilotScope

- Research framework focused on ML4DB
- Not limited to query optimization
 - E.g., also support for index recommendation
- Developed at Alibaba
- Uses Push/Pull connectors on an actual DBMS

<https://github.com/alibaba/pilotscope>



	Broad support for different optimization approaches
	Transparent implementation of prototypes
	Execution on Postgres
	Easy comparison of different prototypes

PostBOUND

Overview

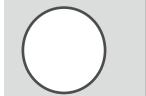
- Python framework for research in query optimization
- Two primary use cases:
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



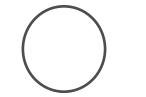
Broad support for different optimization approaches



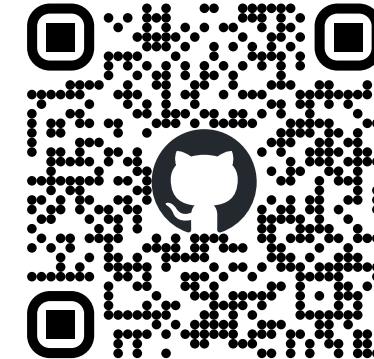
Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes





Overview

- Python framework for research in query optimization
- Two primary use cases
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



Broad support for different optimization approaches



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

- Provides different optimization pipelines
- Pipelines correspond to different optimizer architectures
- Researchers select suitable pipeline
- Researchers implement interfaces of the pipeline



Overview

- Python framework for research in query optimization
- Two primary use cases
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



Broad support for different optimization approaches



Transparent implementation of prototypes



Query execution on Postgres



Easy comparison of different prototypes

- Use hinting to enforce PostBOUND's optimization decisions
- Hints overwrite optimizer of the target database



Overview

- Python framework for research in query optimization
- Two primary use cases
 - Prototyping of new optimization strategies
 - Transparent evaluation of different optimizers



Broad support for different optimization approaches



Transparent implementation of prototypes



Query execution on Postgres



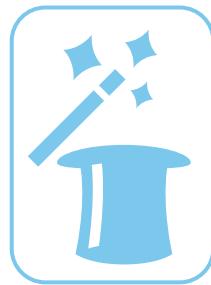
Easy comparison of different prototypes

- Benchmark “support suite”, e.g.,
- Query plan analysis
- Automatic cache prewarming
- Pre-defined workloads + data sets (IMDB, Stats, Stack, SSB)

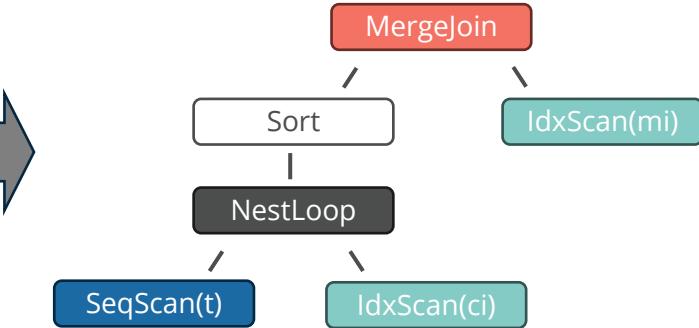
Plan Enforcement via Query Hinting



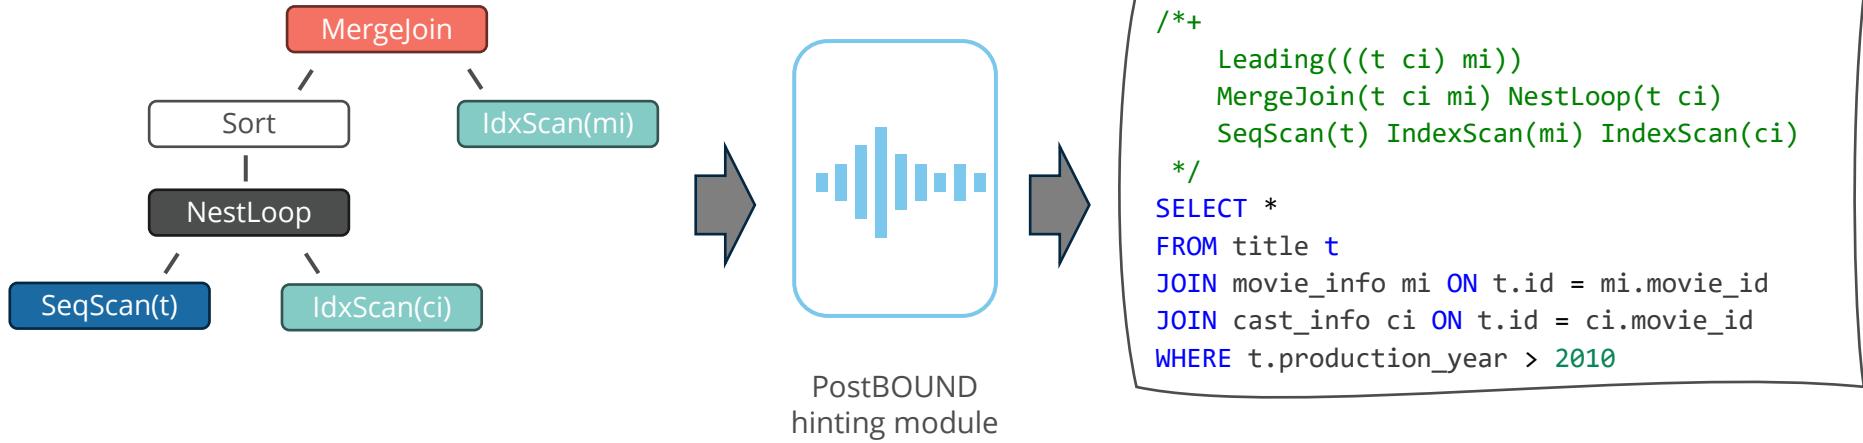
```
SELECT *  
FROM title t  
JOIN movie_info mi ON t.id = mi.movie_id  
JOIN cast_info ci ON t.id = ci.movie_id  
WHERE t.production_year > 2010
```



PostBOUND “magic”
(for now)



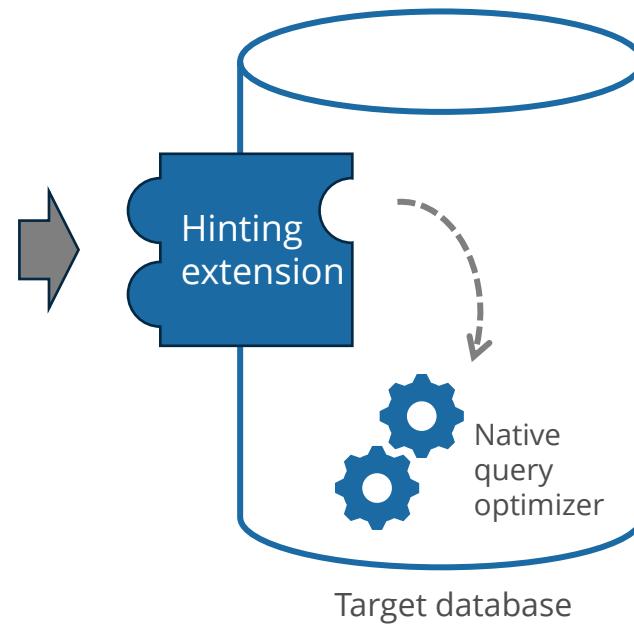
Plan Enforcement via Query Hinting



Plan Enforcement via Query Hinting



```
/*+
  Leading(((t ci) mi))
  MergeJoin(t ci mi) NestLoop(t ci)
  SeqScan(t) IndexScan(mi) IndexScan(ci)
*/
SELECT *
FROM title t
JOIN movie_info mi ON t.id = mi.movie_id
JOIN cast_info ci ON t.id = ci.movie_id
WHERE t.production_year > 2010
```

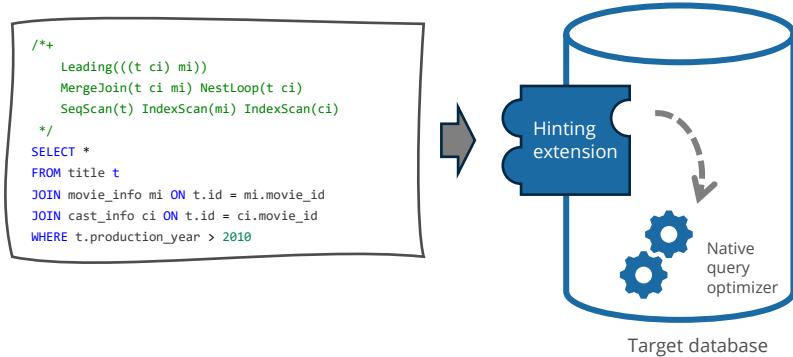


Plan Enforcement via Query Hinting



Supported Hinting Backends

- pg_hint_plan
 - Hints for join order
 - Hints for most physical operators
 - Hints for join cardinalities
- pg_lab
 - Same as pg_hint_plan, but additionally:
 - Hints for parallel plans
 - Hints for base table cardinalities



Consequences

- Can only enforce plan features that are supported by the hinting backend
- Currently no support for re-optimization approaches

"Postgres for Optimizer Research"

- Lightweight fork of PostgreSQL
- Adds additional extension points to the optimization workflow
- Provides fine-grained hinting capabilities

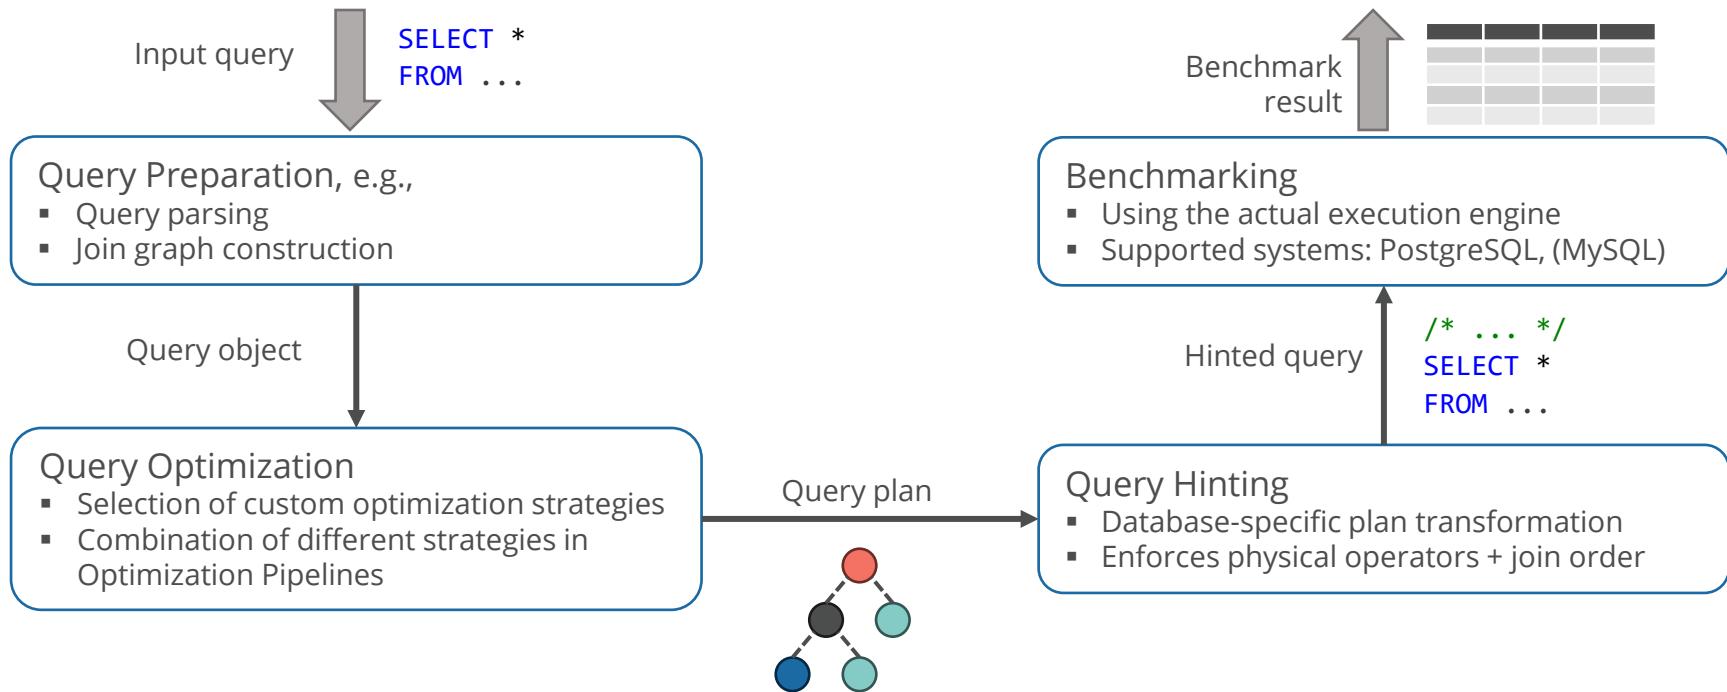


Extension points

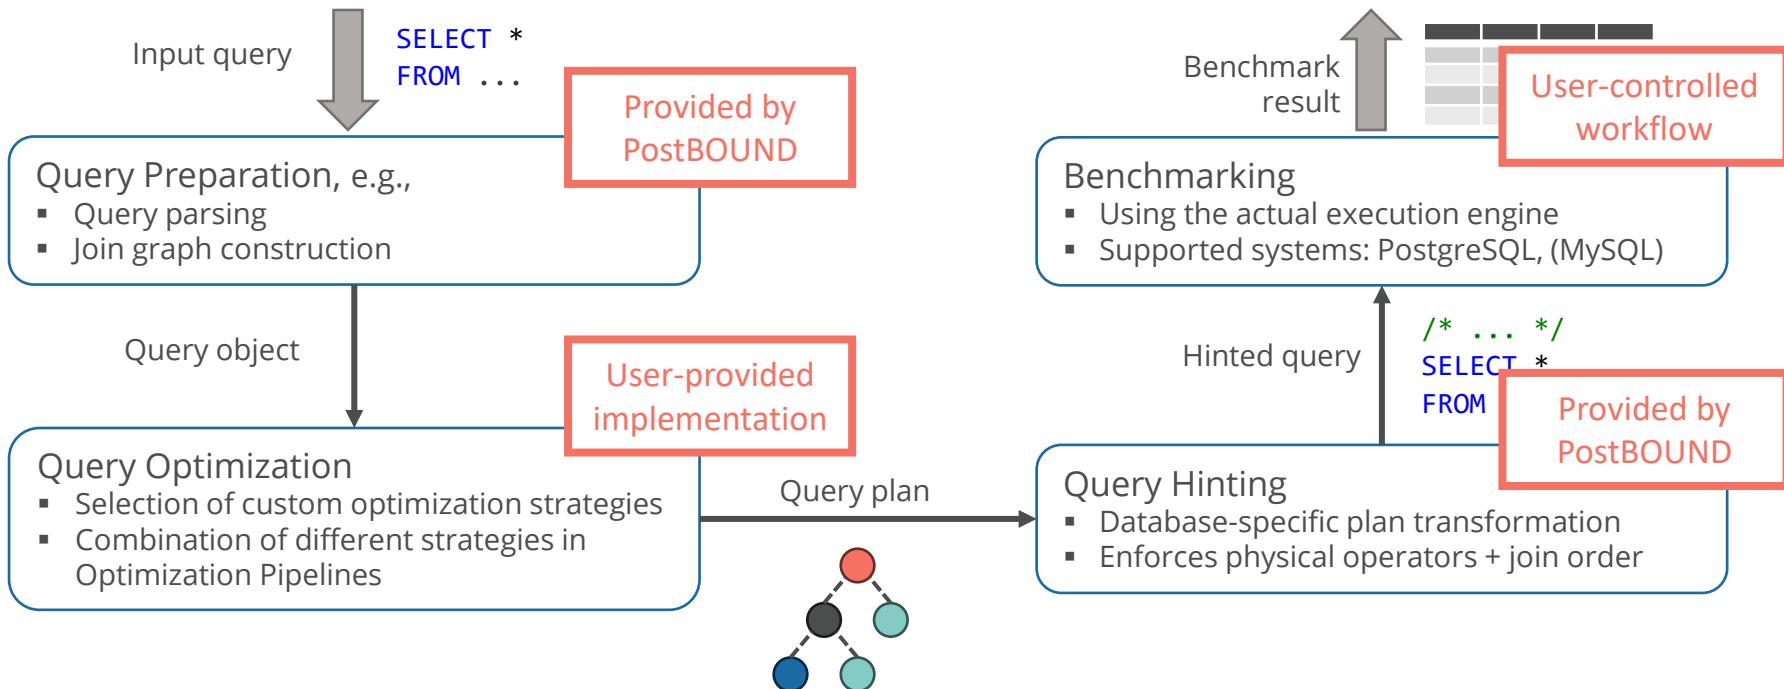
- Cardinality estimation
- Cost estimation
- Path pruning
- Parallel worker estimation



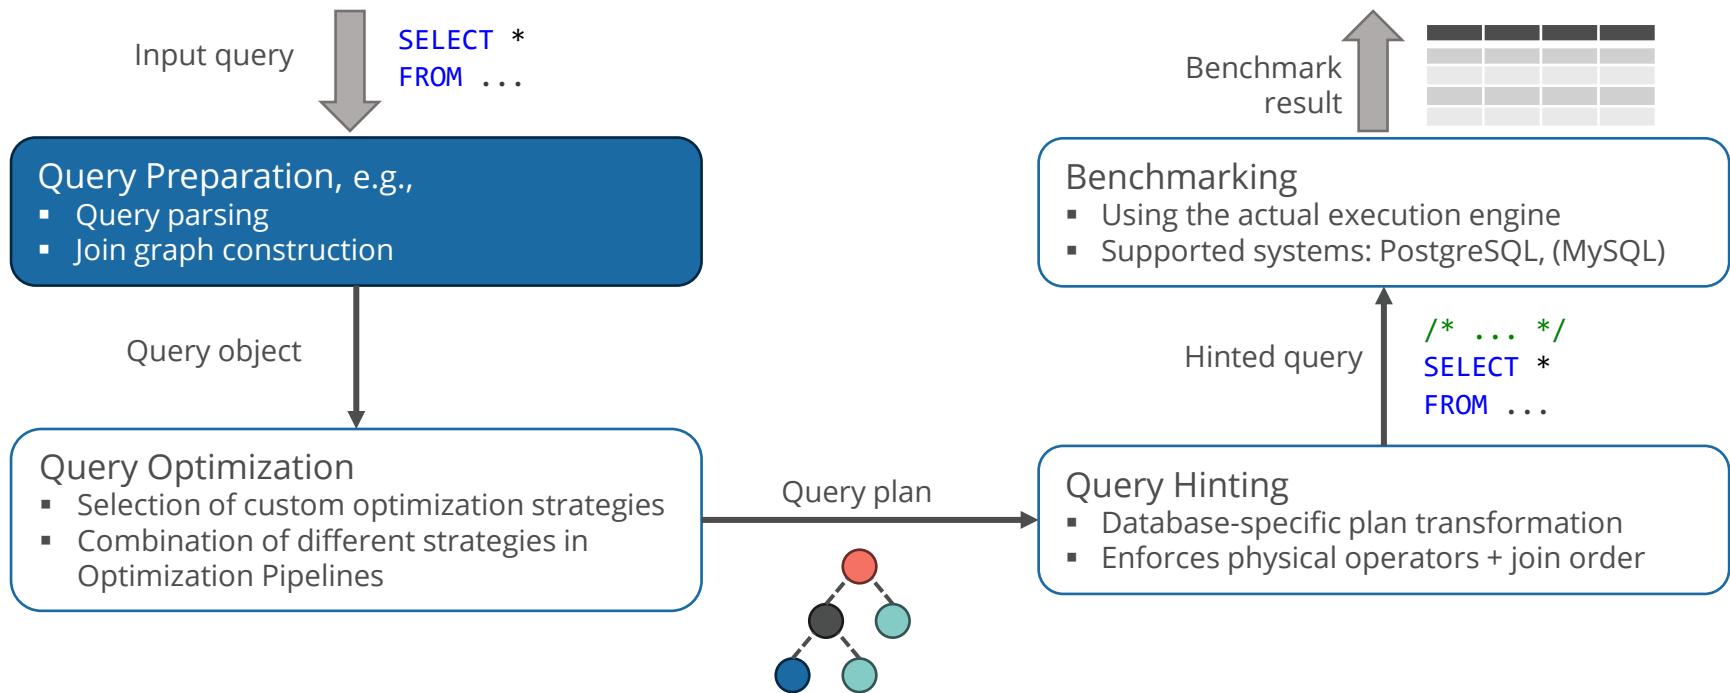
PostBOUND Workflow



PostBOUND Workflow



PostBOUND Workflow



Stats Benchmark

Dataset

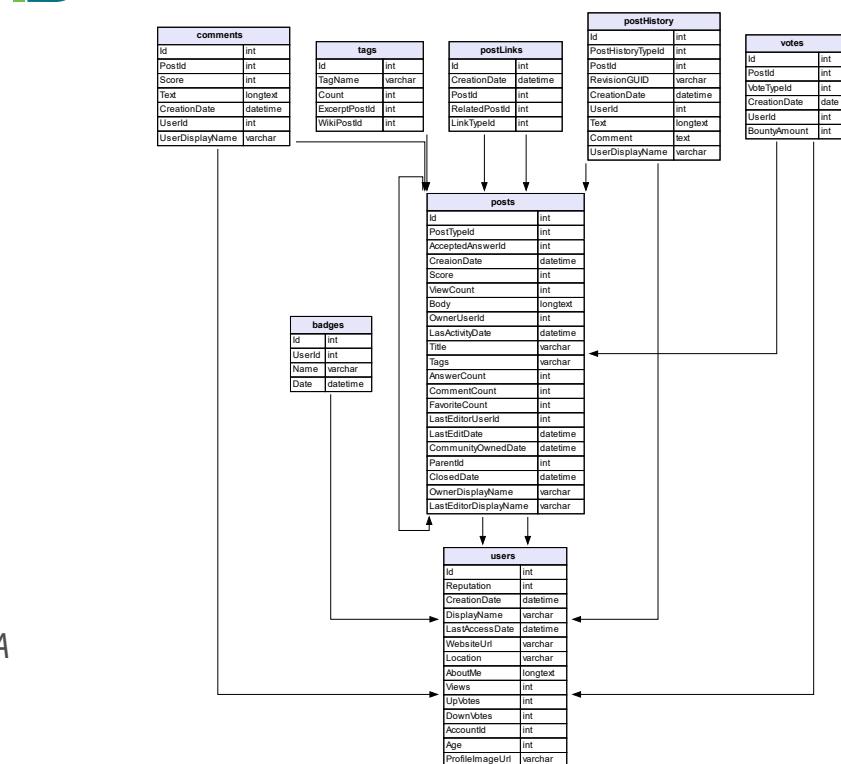
- Based on the Statistics Stackoverflow ([CrossValidated](#))
- Models relation between users, their questions, and answers
- 8 tables with complex (foreign-key) correlations

Workload

- 146 queries in total, between 1 and 6 joins
- Auto-generated queries based on complex templates

Purpose

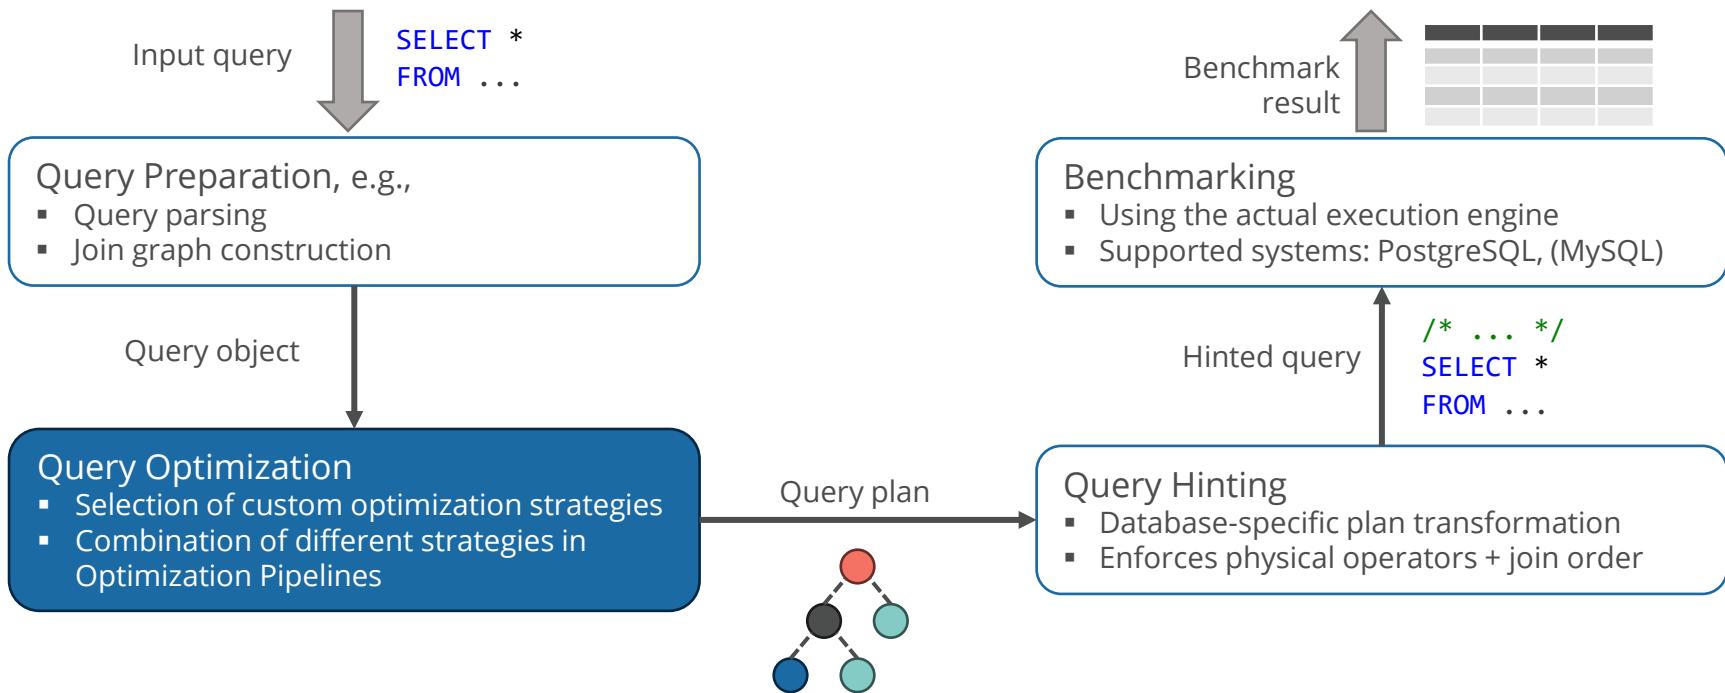
- Lightweight benchmark (<125MB compressed) for cardinality estimation
- Challenging correlations, skewed data distributions
- Presented by Han et al.: “*Cardinality Estimation in DBMS: A Comprehensive Benchmark Evaluation*” (VLDB 2022, [Link](#))





Demo

PostBOUND Workflow



Optimizer Implementation

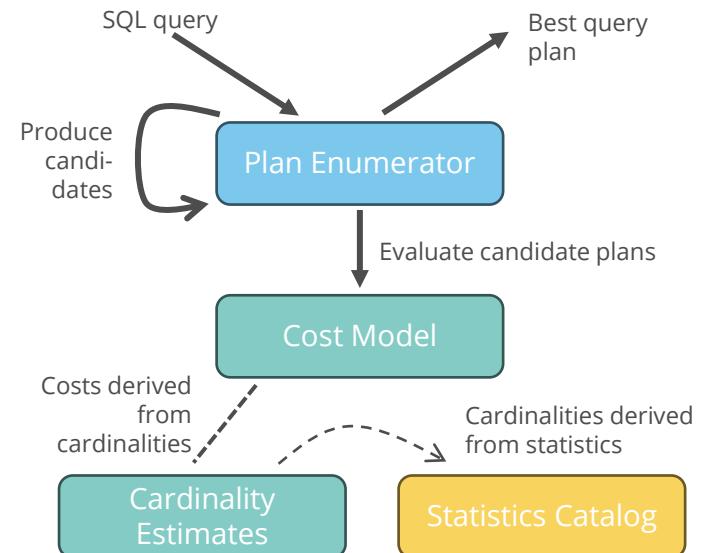


Optimization pipelines

- Mental models for different optimizer architectures
- Pipelines define specific optimization steps (e.g., cardinality estimation)
- Researchers need to implement (subsets of) the stages
- “Fill the gaps” principle
 - Researchers only need to worry about their specific stages
 - PostBOUND uses reasonable defaults for the rest

Textbook pipeline

- Available stages:
 - Plan enumerator
 - Cost model
 - Cardinality estimator
- Defaults: native estimators and simulated enumerator



Optimizer Implementation

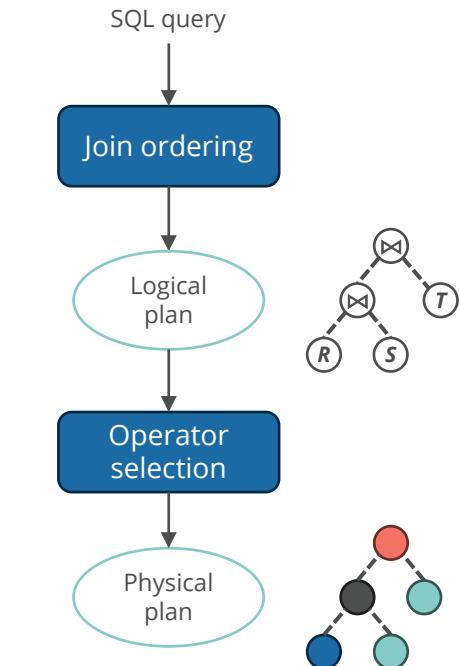


Optimization pipelines

- Mental models for different optimizer architectures
- Pipelines define specific optimization steps (e.g., cardinality estimation)
- Researchers need to implement (subsets of) the stages
- “Fill the gaps” principle
 - Researchers only need to worry about their specific stages
 - PostBOUND uses reasonable defaults for the rest

Multi-stage pipeline

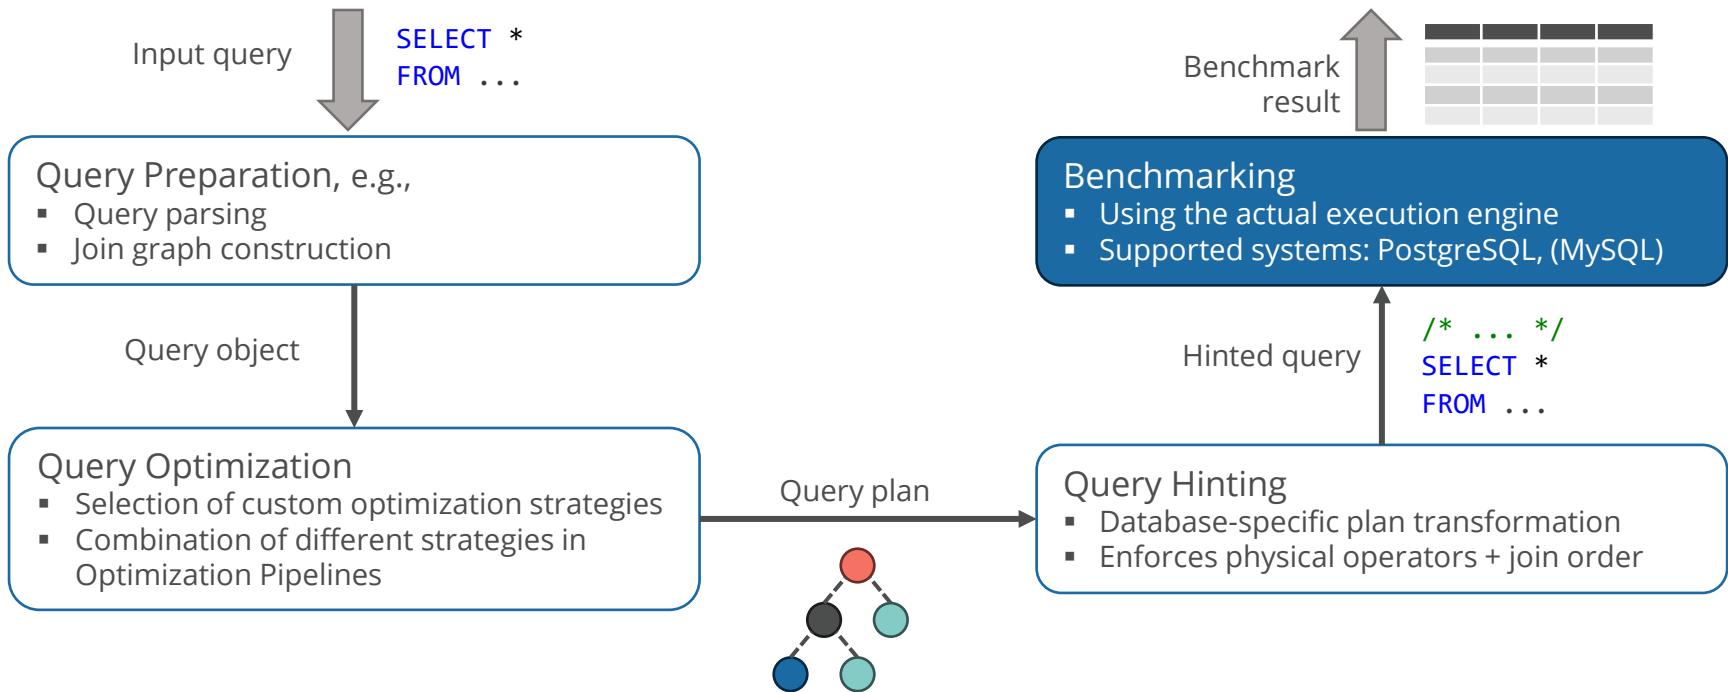
- Available stages:
 - Join ordering
 - Operator selection
 - (Plan parameterization)
- Defaults: native strategies





Demo

PostBOUND Workflow



Benchmarking



Benchmark execution

- Goal: reproducible + transparent evaluation
- Query preparation for common tasks (e.g., warm start)
- Tools generate many additional data points
 - Pipeline configuration
 - Database schema
 - ...
- Integrated into Python data analysis ecosystem (e.g., pandas)



Support Tooling

- Setup scripts for common databases + workloads (IMDb, SSB, Stats, StackOverflow)
- Setup scripts for different Postgres versions





Demo



Implementing Your Own Optimizer

Hands-On Material:
github.com/db-tu-dresden/SIGMOD25-OptimizerTutorial



Supervised Cardinality Estimation (*MSCN-light*)



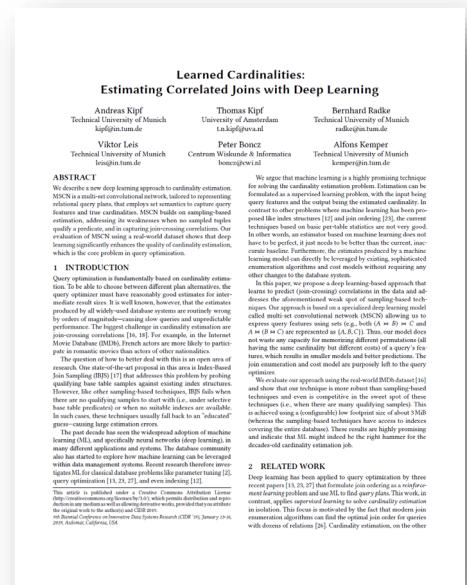
MSCN in a nutshell

- Supervised cardinality estimation model
 - Queries are featurized according to
 - Their tables
 - The join graph
 - Filter predicates
 - Different sub-modules for each component
 - Modules are merged in final model

Our adaptation

- Feature each component with an embedding model
 - Predict cardinality with a simple gradient boosting model

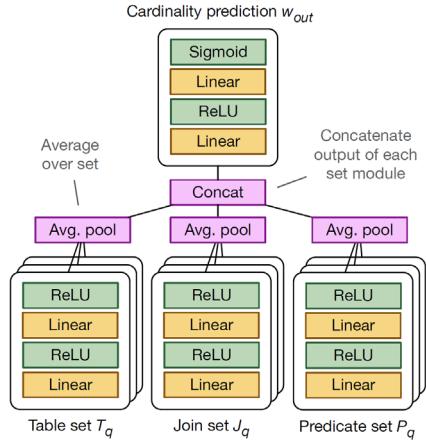
Kipf et al.: "Learned Cardinalities: Estimating Correlated Joins with Deep Learning" (CIDR'2019)



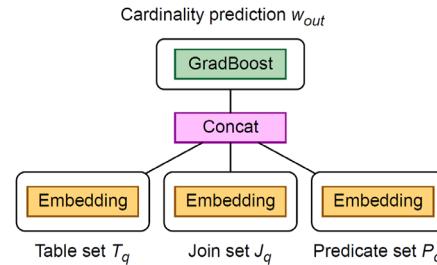
Supervised Cardinality Estimation (*MSCN-light*)



Original MSCN model



Light model



Kipf et al.: "Learned Cardinalities: Estimating Correlated Joins with Deep Learning" (CIDR'2019)



Live!

Plan Selection (*BAO-light*)



BAO in a nutshell

- Use (global) planner hints to obtain different query plans
- Vectorize query plans using binarization + schema-independent features
- Estimate plan quality using a TCNN
- Selects the (estimated) best plan for execution

Our adaptation

- Use embedding to vectorize query plan
- Predict plan cost with a simple gradient boosting model

Research Data Management Track Paper
SIGMOD '21, June 20–25, 2021, Virtual Event, China

Bao: Making Learned Query Optimization Practical

Ryan Marcus
MIT & Intel Labs
rym@mit.edu
Nesime Tatbul
MIT & Intel Labs
tatbul@csail.mit.edu
Parimkar Negi
MIT
pnegi@mit.edu
Mohammad Alizadeh
MIT
alizadeh@csail.mit.edu
Hongzhi Mao
MIT
hongzhi@mit.edu
Tim Kraska
MIT
kraska@csail.mit.edu

ABSTRACT
Recent efforts applying machine learning techniques to query optimization have shown promising results. However, these approaches often suffer from overhead, inability to adapt to changes, and poor tail performance. In this paper, we propose Bao, a learned query optimizer that can make probabilistically expensive operations in practice (thus why we wish to call it “practical”) more efficient. Bao takes advantage of the wisdom built into existing query optimizers by providing proxy query optimization hints. Bao uses a novel sampling, a well-studied reinforcement learning algorithm. As opposed to previous reinforcement learning approaches, Bao does not require changes in query workload, data, and schema. Experimentally, we show that Bao can outperform traditional optimizers on a wide range of end-to-end query execution performance, including tail latency for several workloads containing long-running queries. In short, Bao achieves better performance compared with a commercial system.

CCS CONCEPTS
• Information systems → Query optimization.

KEYWORDS
• Databases, machine learning, reinforcement learning

ACM Reference Format:
Ryan Marcus, Parimkar Negi, Hongzhi Mao, Nesime Tatbul, Mohammad Alizadeh, and Tim Kraska. 2021. Bao: Making Learned Query Optimization Practical. In Proceedings of the 2021 International Conference on Management of Data (SIGMOD '21), June 20–25, 2021, Virtual Event, China, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3441299.3452038>

1 INTRODUCTION
Query optimization is an important task for database management systems (DBMSs). It is a multi-step process that involves elements of query optimization – cardinality estimation and cost modeling – to produce efficient query execution plans [38]. Several works have applied machine learning (ML) to query optimization [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54]. While all of these new solutions demonstrate significant improvements, they also highlight that some of these techniques are yet practical, as they suffer from several fundamental problems:

- (1) Long training time. Most proposed machine learning techniques require a large amount of training data to learn, which may have a positive impact on query performance. For example, ML-powered cardinality estimators based on supervised learning require millions of training examples to learn [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54].
- (2) Inability to adjust to data and workload changes. While proposed reinforcement learning approaches can handle unpredictable changes in query workload, data, or schema can make them less effective. For example, if a DBMS undergoes schema evolution, it must be retrained when data changes, or risk becoming stale.

Several proposed reinforcement learning techniques assume that both the query workload and schema remain static during the entire complete retraining when this is not the case [40, 31, 32, 39].

- (3) Tail latency. Reinforcement learning techniques can outperform traditional optimizers on average, but often perform catastrophically (e.g., 10x regression in query performance) when the DBMS undergoes schema evolution between training data to query. While some approaches offer statistical guarantees for the performance of reinforcement learning approaches, even if rare, are unacceptable in many real world applications.
- (4) Lack of explainability. Reinforcement learning approaches are already complex, understanding query optimization is even harder when black box deep learning approaches are used. More research is needed to understand how learned query optimizers do not provide a way for database administrators to influence the query execution plan.
- (5) Integrations costs. To the best of our knowledge, all previous learned optimizers are still research prototypes, offering little to no integration with production DBMSs. This is due to the lack of standard SQL, not to mention vendor specific features. Hence, fully integrating learned optimizers into production DBMSs is a non-trivial task. The database administrator (DBA) just needs to download our open-source module¹, and even has the option to selectively turn the learned optimizer on or off for specific queries.

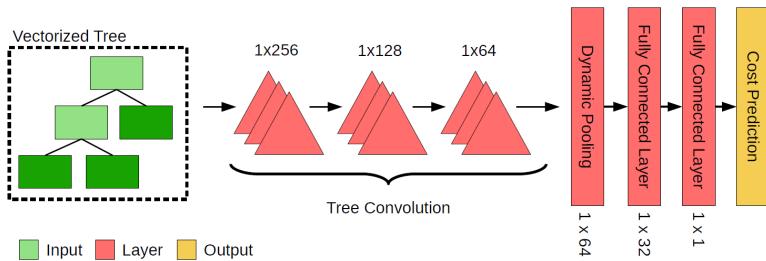
¹ <https://github.com/kraska/Bao>

Marcus et al.: “BAO: Making Learned Query Optimization Practical” (SIGMOD’2021)

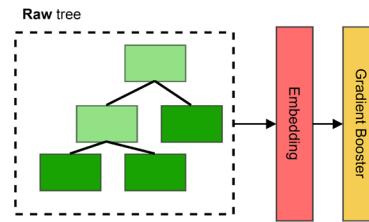
Plan Selection (*BAO-light*)



Original BAO model



Light model



Marcus et al.: "BAO: Making Learned Query Optimization Practical" (SIGMOD'2021)



Live!

Beyond Learned Optimizers: Pessimistic



Core idea

- Use upper bounds instead of cardinality point-estimates to derive a robust execution plan
- Disable dangerous join operators, e.g., nested-loop join

Our adaptation

- Use the UES bound formula
- Only enable hash joins
- DBMS is free to select a join order
- DBMS is free to select scan operators

Hertzschuch et al.: "Simplicity Done Right for Join Ordering" (CIDR'2021)

Simplicity Done Right for Join Ordering

Axel Hertzschuch, Claudio Hartmann, Dirk Habich, Wolfgang Lehner
Database Systems Group
Technische Universität Dresden
Dresden, Germany
firstname.lastname@tu-dresden.de

ABSTRACT
In this paper, we propose a simple, yet fast and effective approach for join ordering selection in query optimizers. Our scheme comprises three building blocks: (i) a simple upper bound for arbitrary multi-join queries, (ii) a sketch-based approach for estimating the cardinality bound, and (iii) sampling as query execution to provide tight bounds for the cardinality of join results. As we are going to show, using the Join-Order-Benchmark, our approach is able to find a good join order for a query with significantly less optimization overhead, resulting in a query execution time of 11.3 ms for all 111 JOIN queries compared to state-of-the-art and recent approaches.

1. INTRODUCTION
Although query optimization has been a core research topic for decades, it is still far from being solved [7]. This is mainly due to the fact that there are many degrees of freedom in query optimization: finding a good join order [2, 7], to choose the right scan order [1, 10], to decide which filters are required [2, 7]. This includes joins over intermediate relations, which is a well-known challenge in query optimization. A natural question whether it is even possible to achieve such estimates without join execution is yet to be answered [7].
Recent work on join cardinality estimation has proposed heuristics that may assume predicates independently and a few others have proposed joint heuristics [1, 10]. Upper bounds for the join cardinality estimation have been proposed in [8, 11, 21] and [2, 7]. The first two approaches are appealing at first glance, for example [8, 11, 21], but they do not scale well to many joins [3, 21]. On the other hand, the third approach [2, 7] is very efficient and scales well [10]. Moreover, we compare our concept with the sketch-based upper bound approach of Cai et al. [2]. As we are going to show, our approach is similar to the sketch-based approach in terms of execution time and significantly less planning time due to the simplicity of our concept compared to the sketch-based approach [2].

2. U-BLOCKS: JOIN ORDERING
This article describes our two fundamental building blocks *U* and *F* for a reliable join ordering. For that, we assume materialized views and a set of base tables (e.g., heap, log, key statistics) and precise filter selectivity estimates. While histograms may suffice for base filters, Section 4 shows that histograms are not sufficient for join filters (S-Blocks).

2.1. U-Block: Simple Upper Bound for Joins
The first building block *U* includes a simple upper bound for the cardinality of a join. To calculate this upper bound, we start with a single join and disease arbitrary joins. We then consider all possible ways of joining the (pre-filtered) tables, we calculate the smallest number of distinct values selected by each join and multiply them.

All materials to reproduce and further analyze the results reported in this paper are available at: <https://gitlab.com/ahertz/SimplicityDoneRight>.



Live!

Reproducible Prototyping of Query Optimizer Components

Rico Bergmann and Dirk Habich

Technische Universität Dresden, Germany

SIGMOD'25 Tutorial, June 27, 2025, Berlin, Germany